- Launch VM instances through AWS
- Allow network traffic (ping, ssh)
- Transfer files
- Download, install, and run Hadoop

# Launch VM instances

Submit your application to Amazon

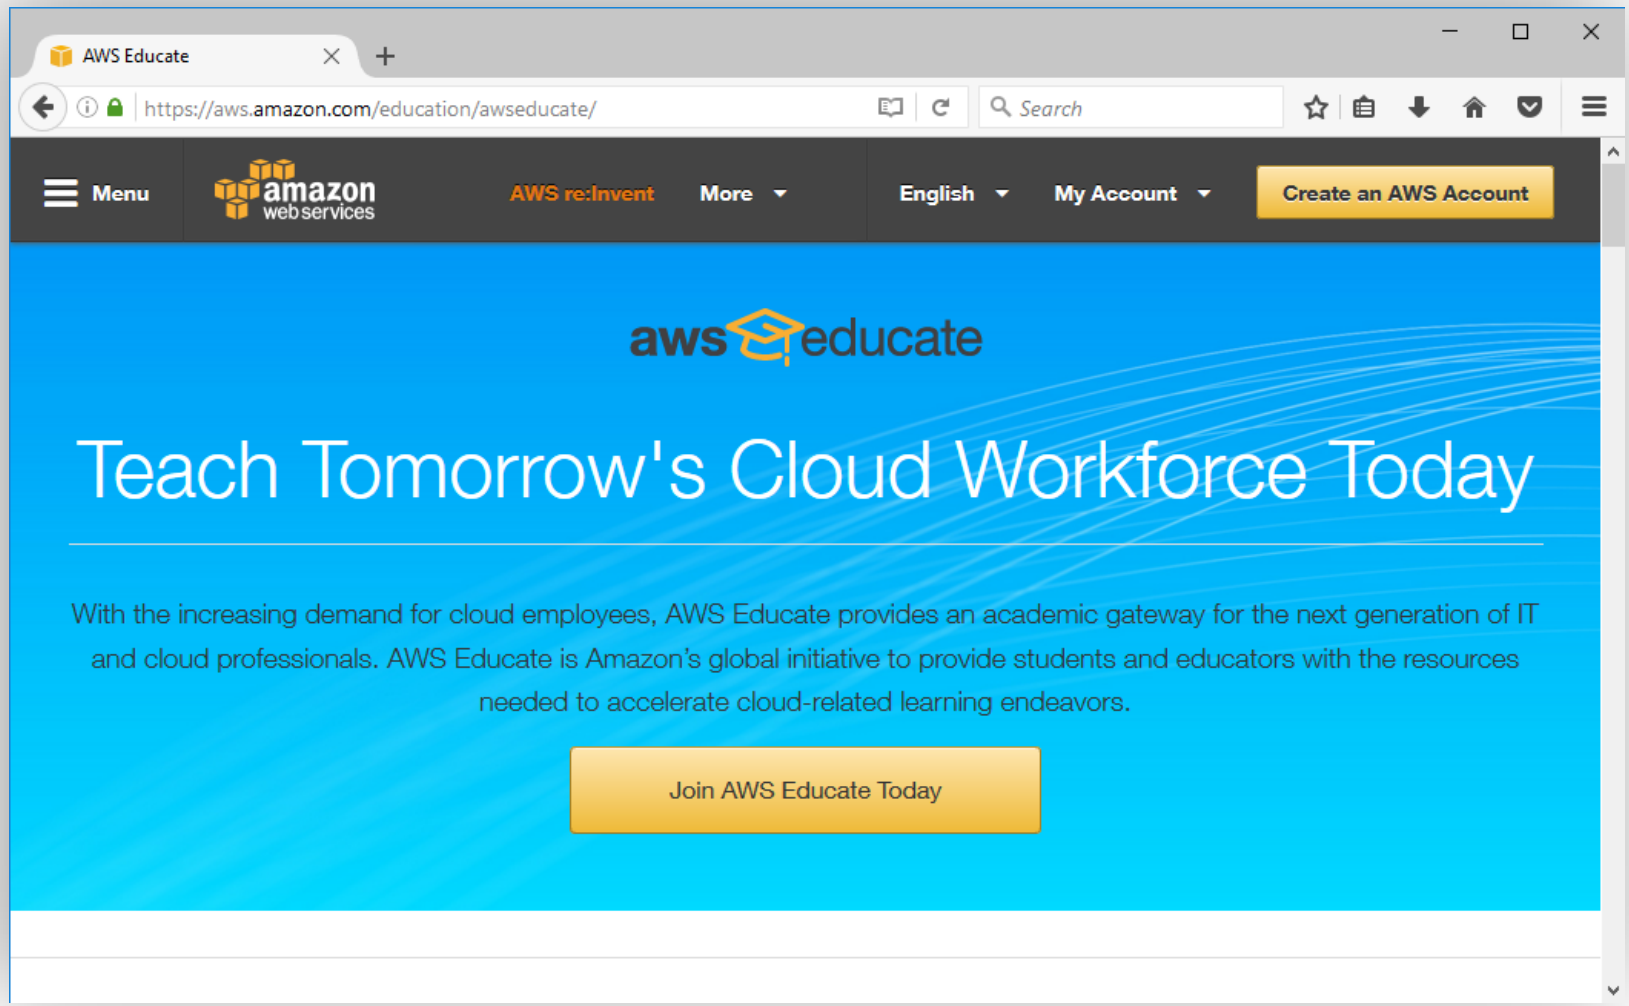

# Launch VM instances

Launch a virtual machine

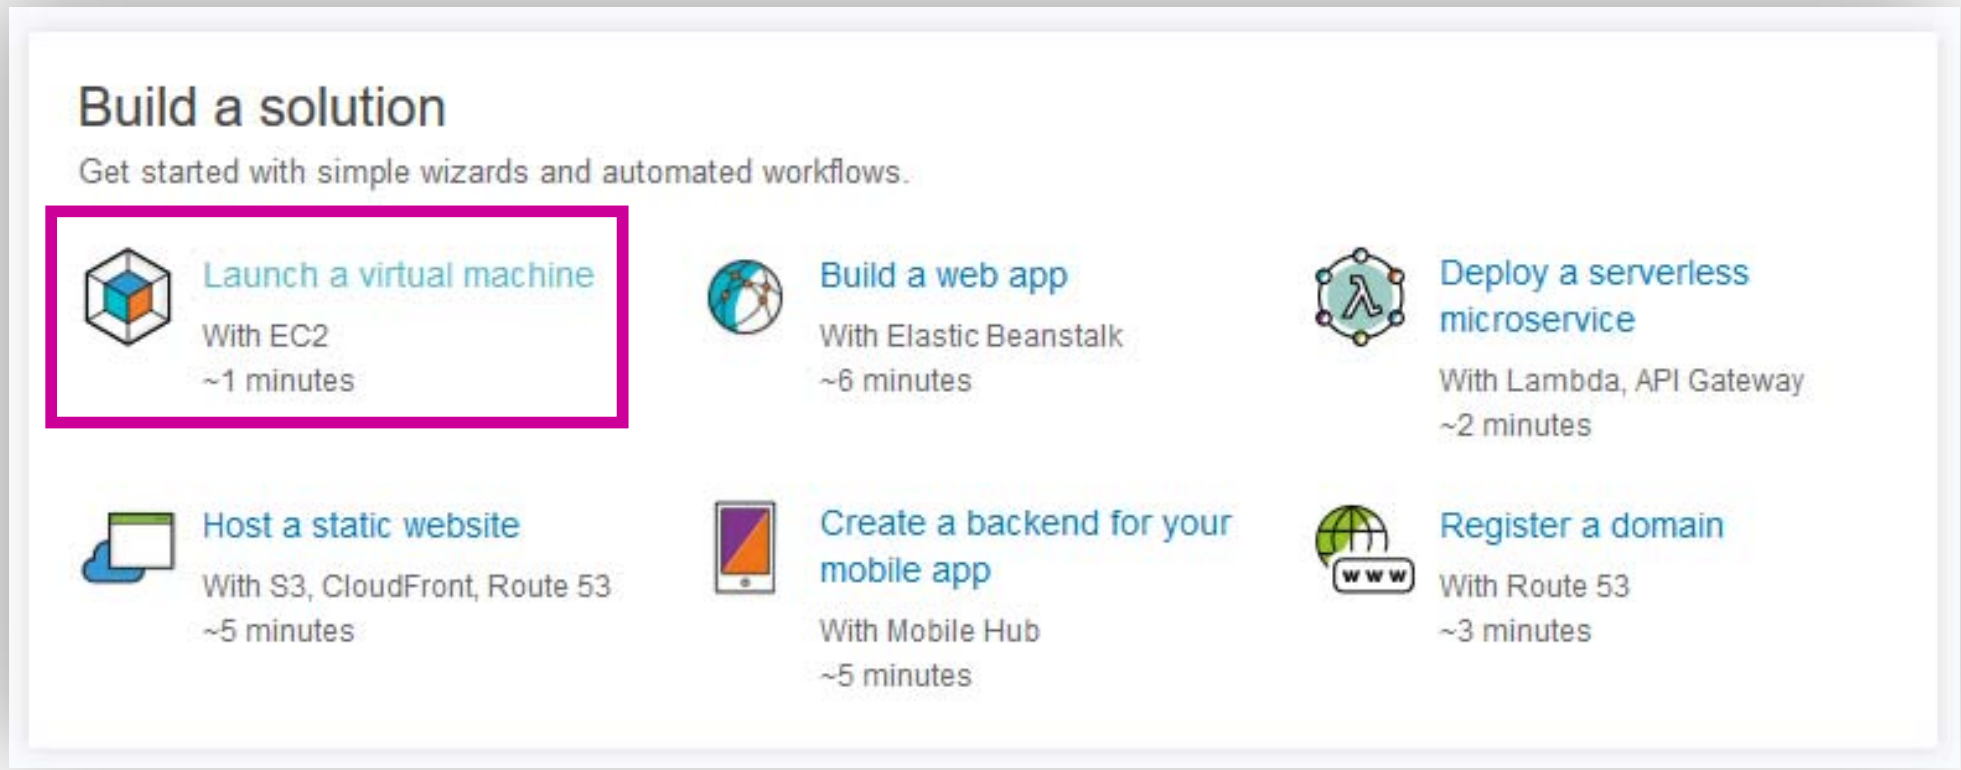

# Launch VM instances

Provide necessary information
Create and download the key pair

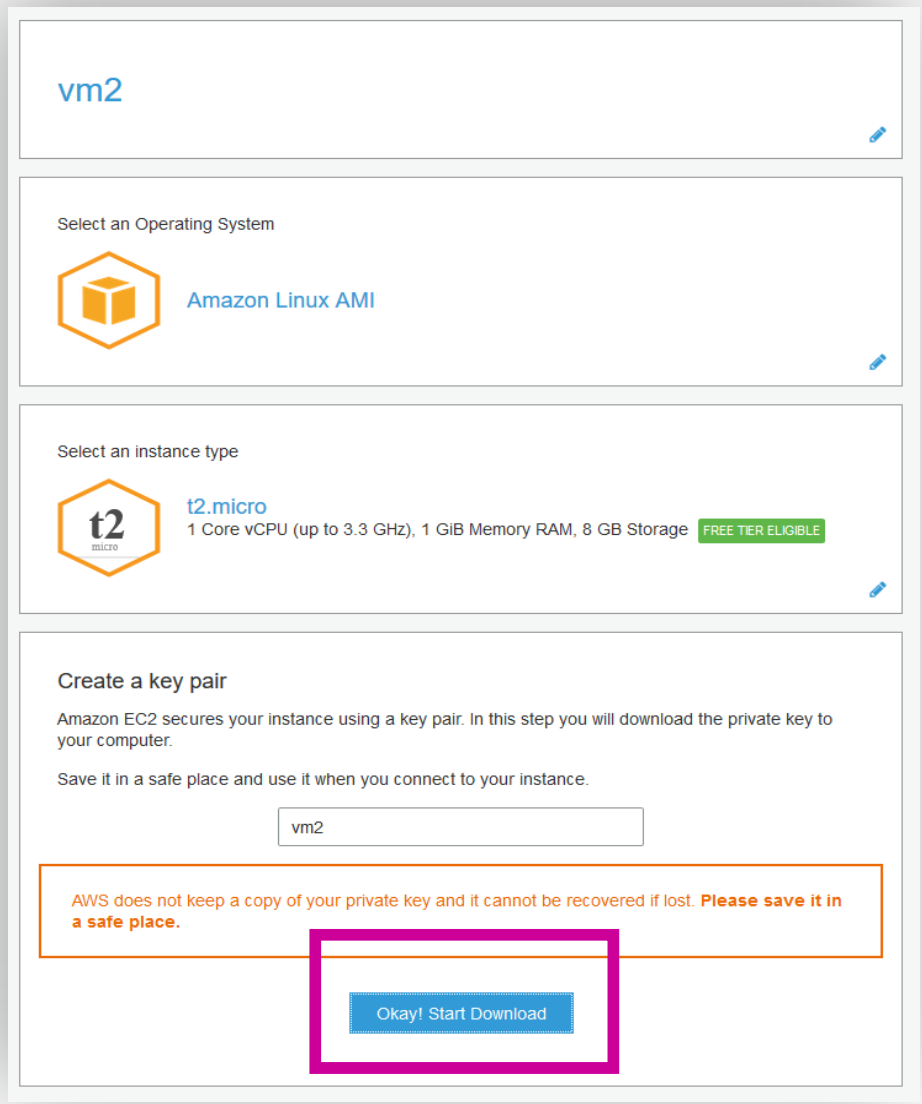

# Launch VM instances

Create the instance

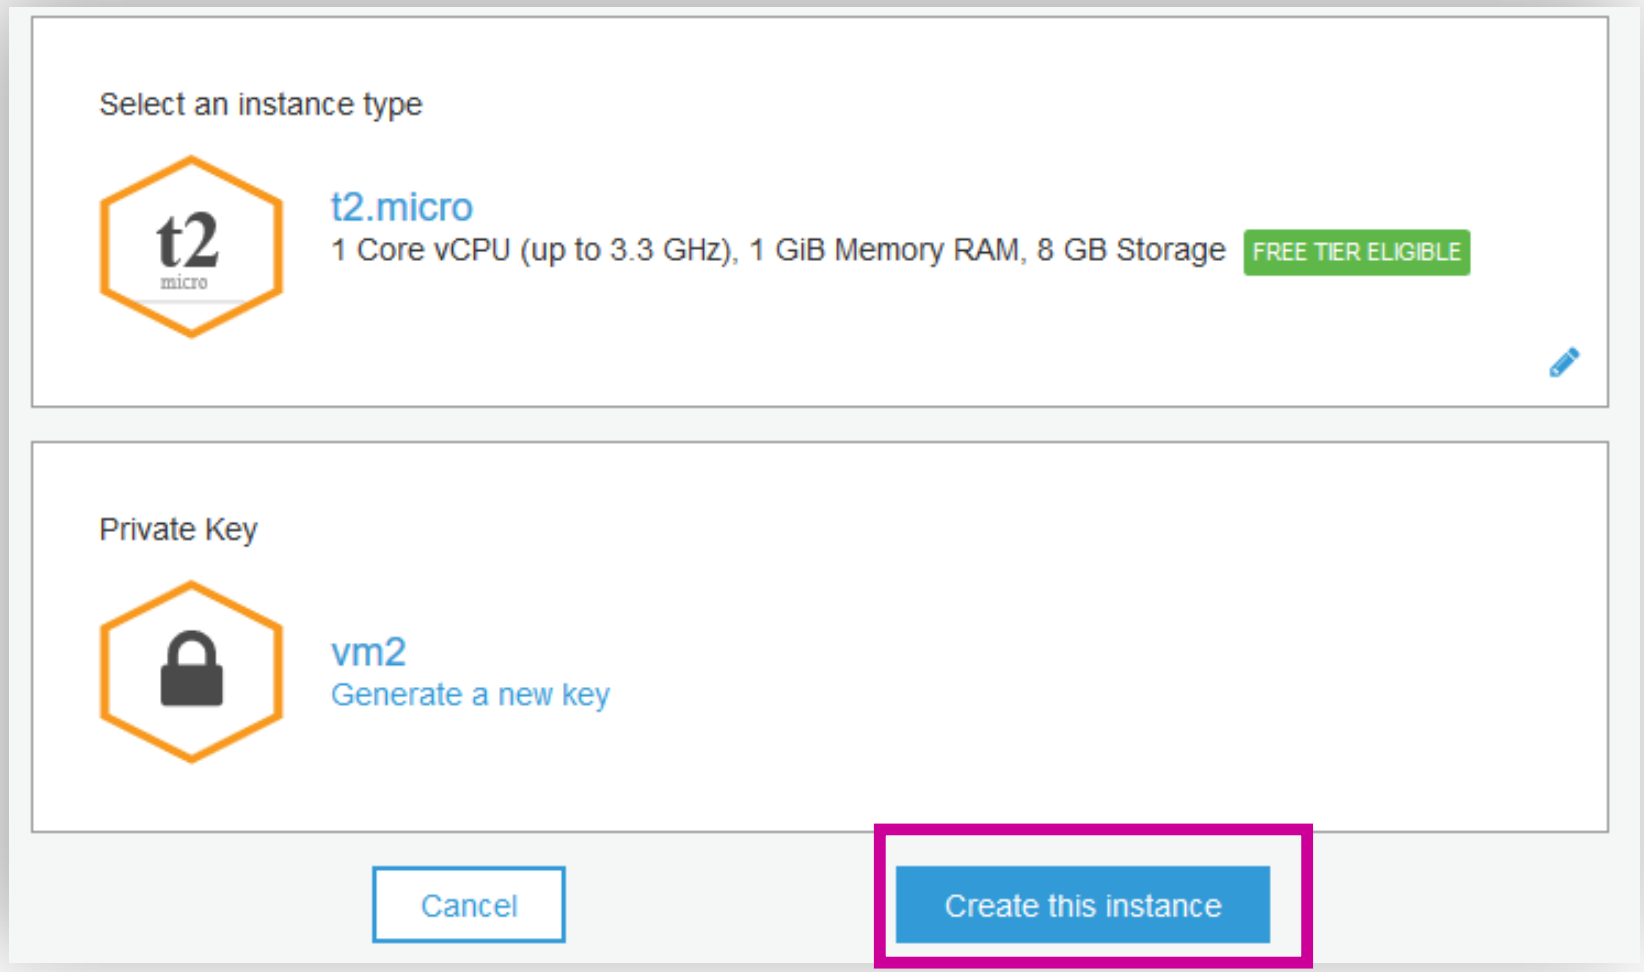

# Launch VM instances

Creating…
You can go ahead proceed to the console

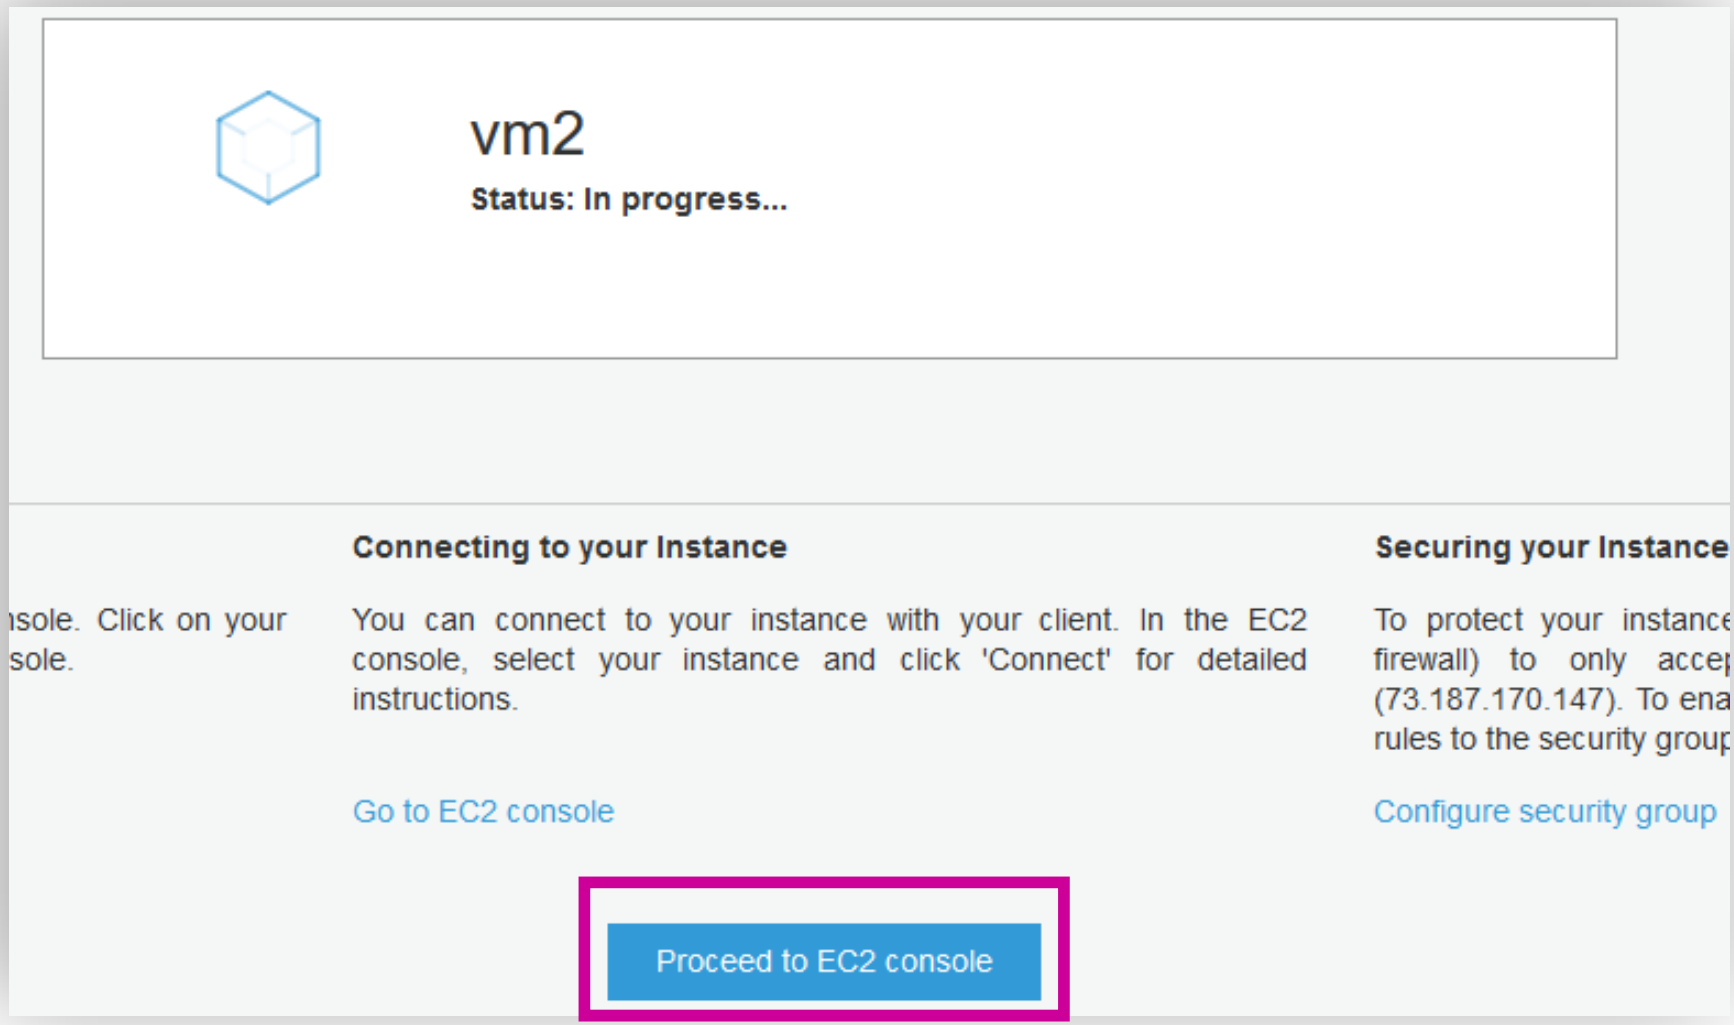

# Launch VM instances

The console is like…

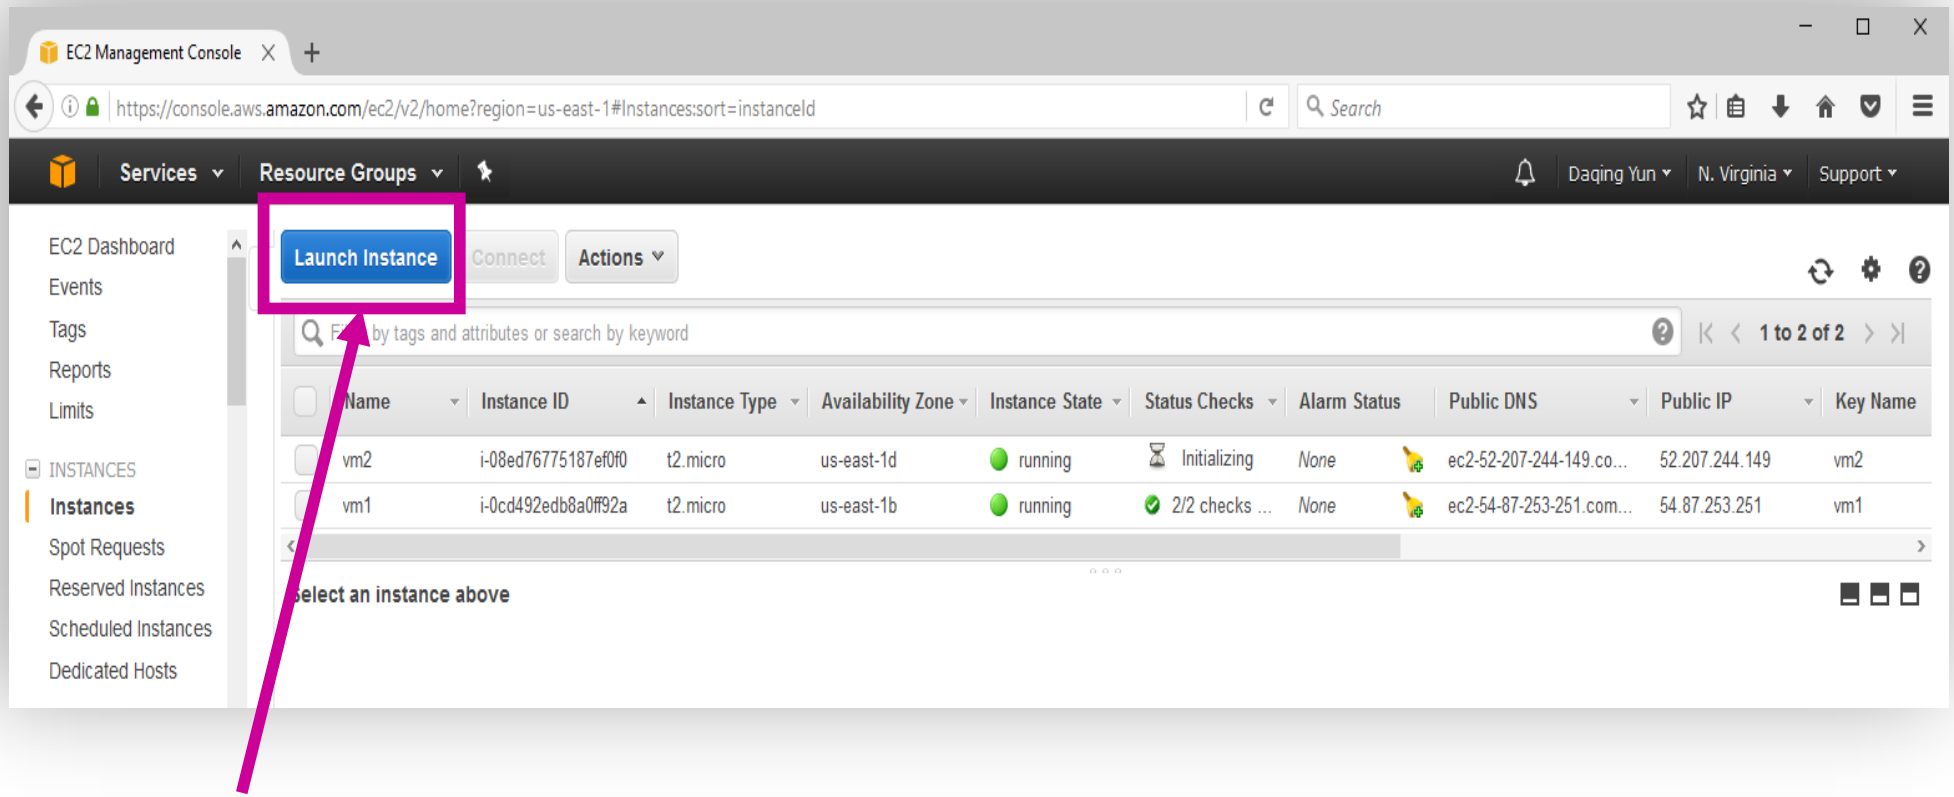


You can keep launching…

# Allow Network Traffic

Try to connect…

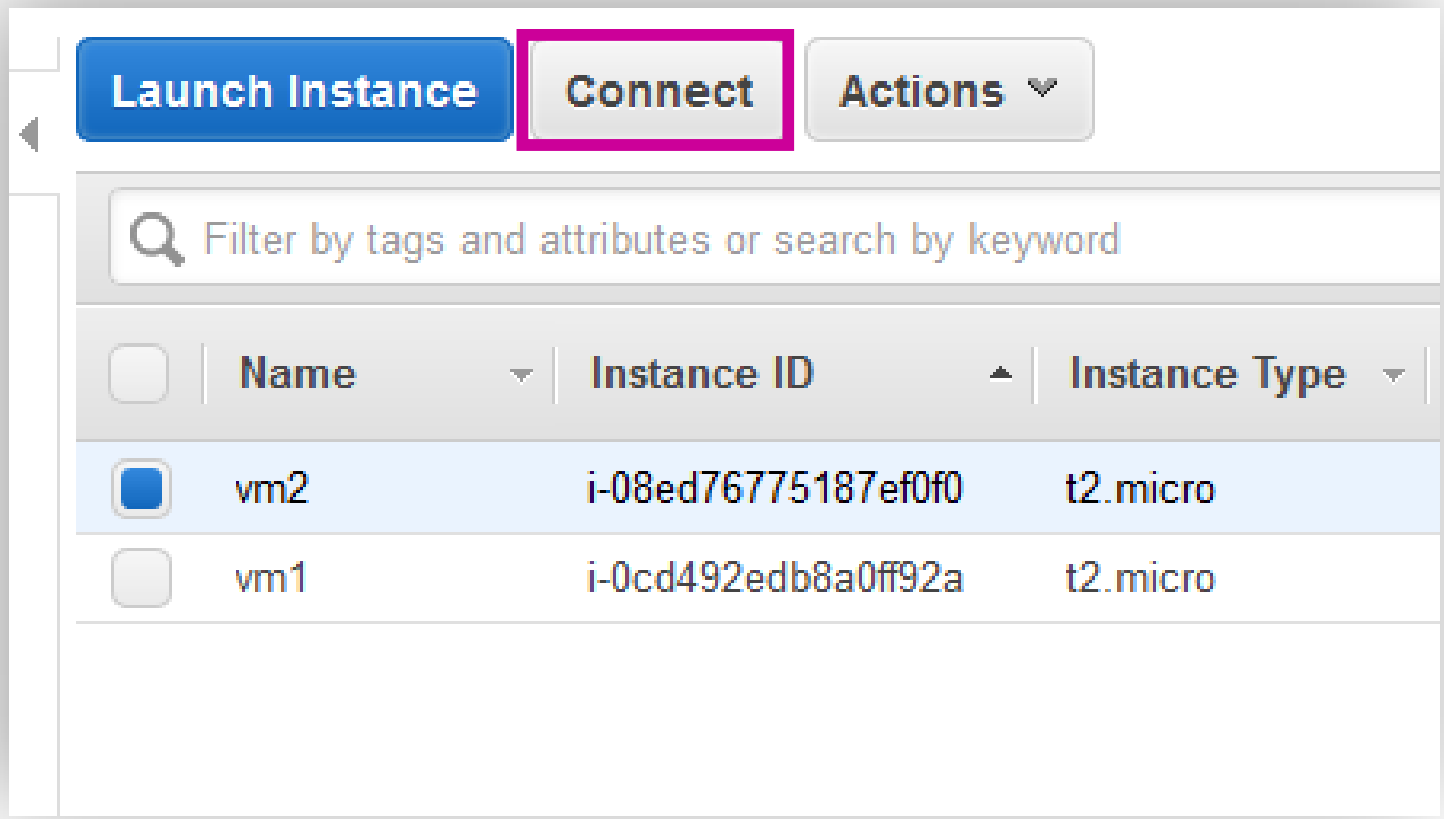

# Allow Network Traffic

Once you click connect, you will see instructions. If you are using Mac, you can use Terminal in /Application to connect

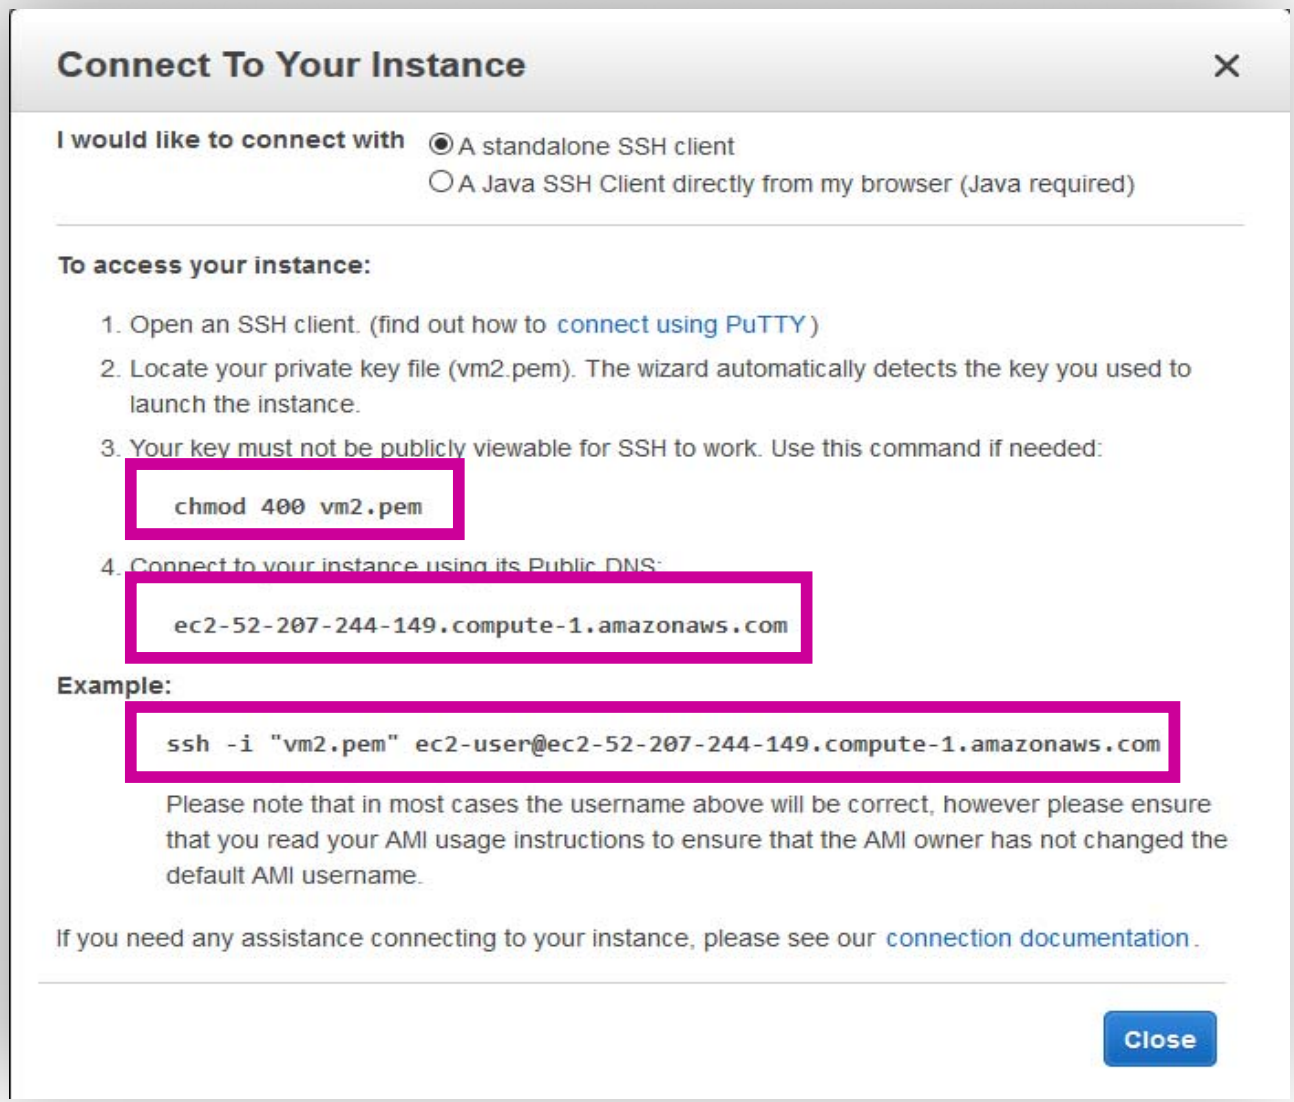

# Allow Network Traffic

Let us use putty to connect from Windows 10

Download and install putty.exe
https://the.earth.li/~sgtatham/putty/latest/x86/putty.exe

Also download puttygen.exe
https://the.earth.li/~sgtatham/putty/latest/x86/puttygen.exe

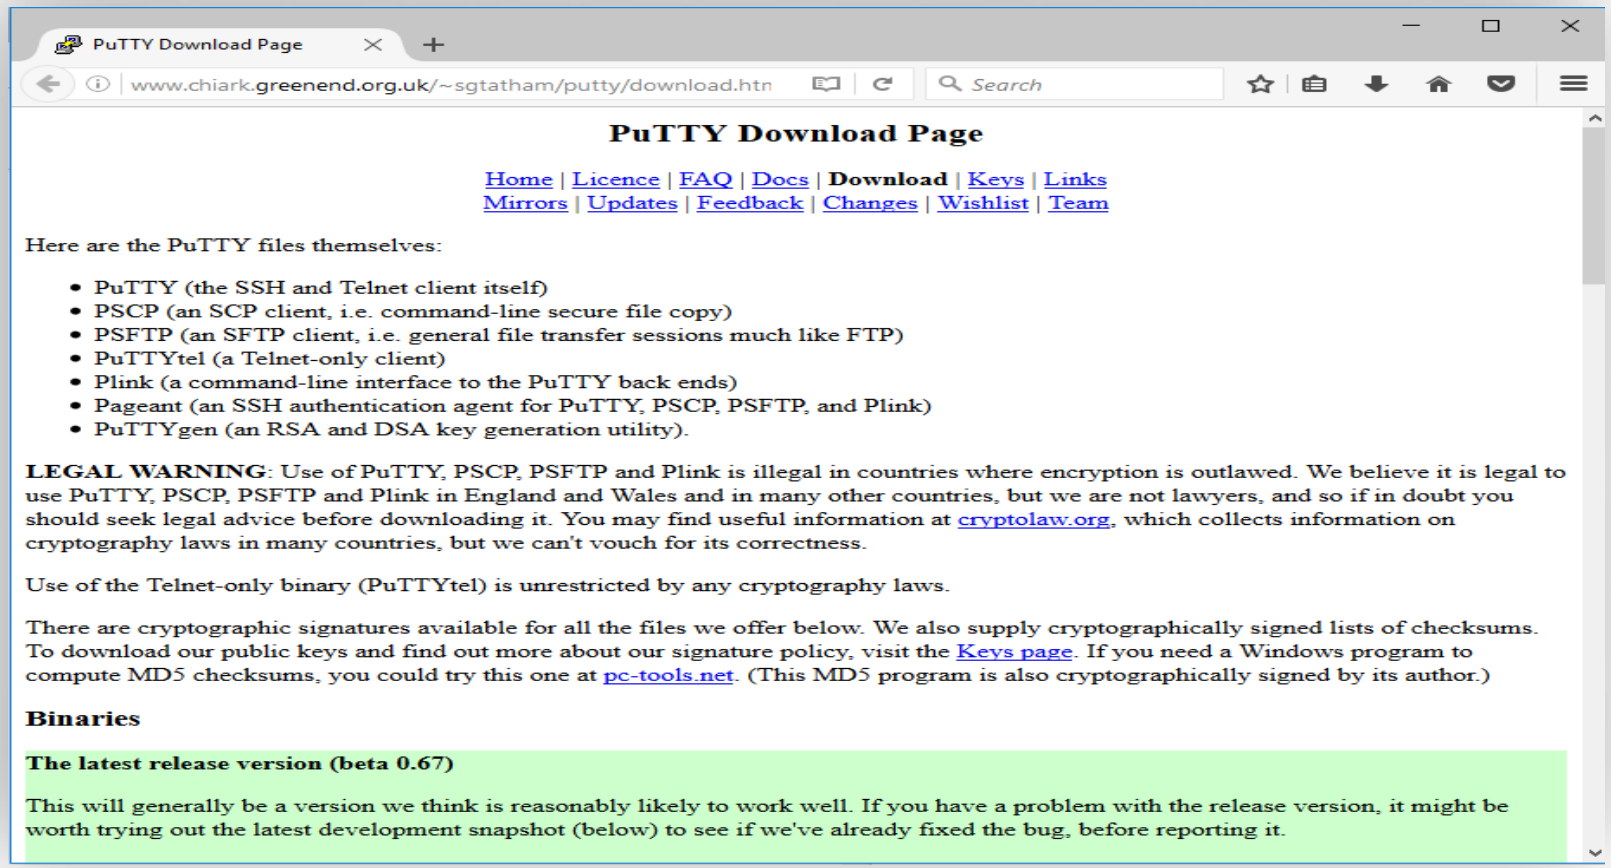

# Allow Network Traffic

PuTTY does not natively support the private key format (.pem) generated by Amazon EC2. Let's convert it using puttygen.exe

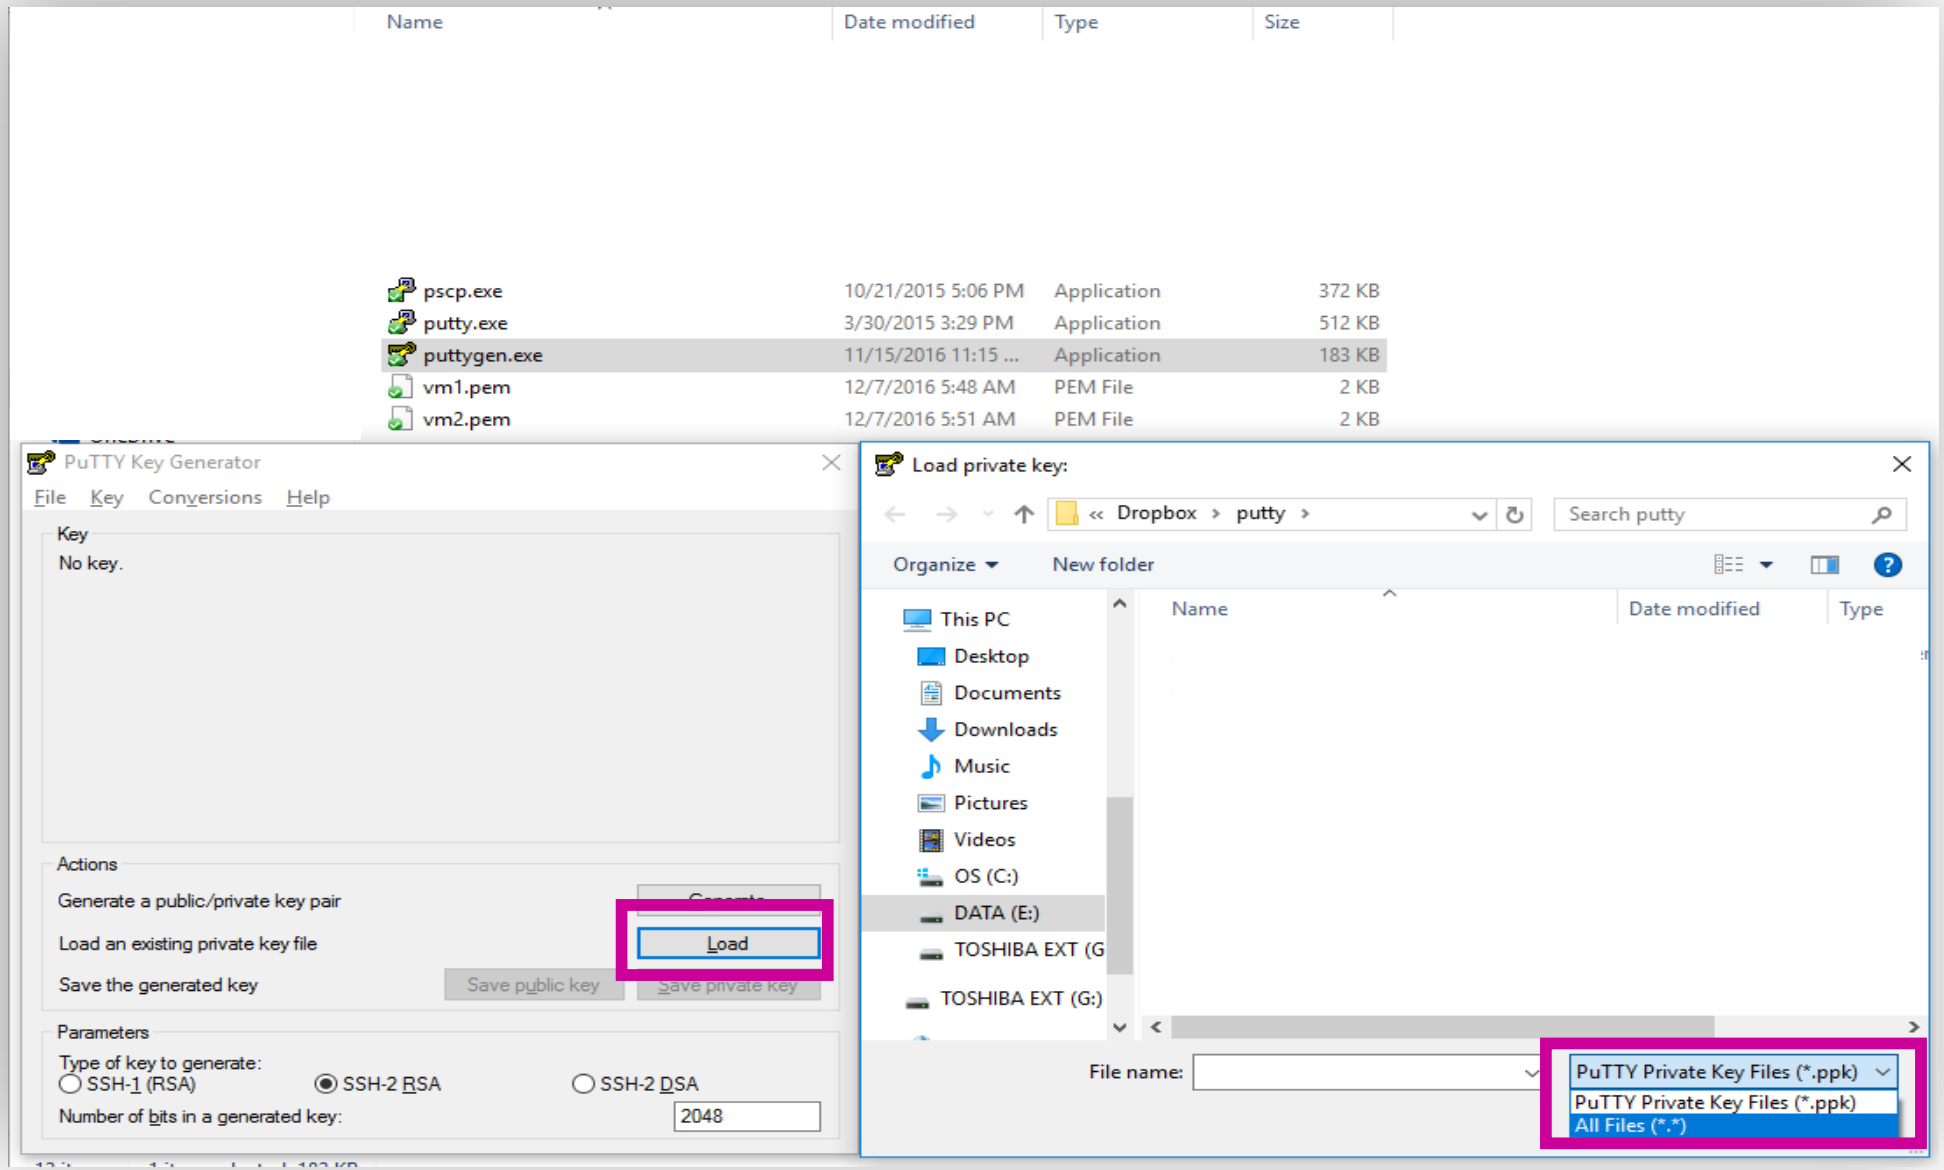

# Allow Network Traffic

Save private key, specify the file name with .ppk file extension

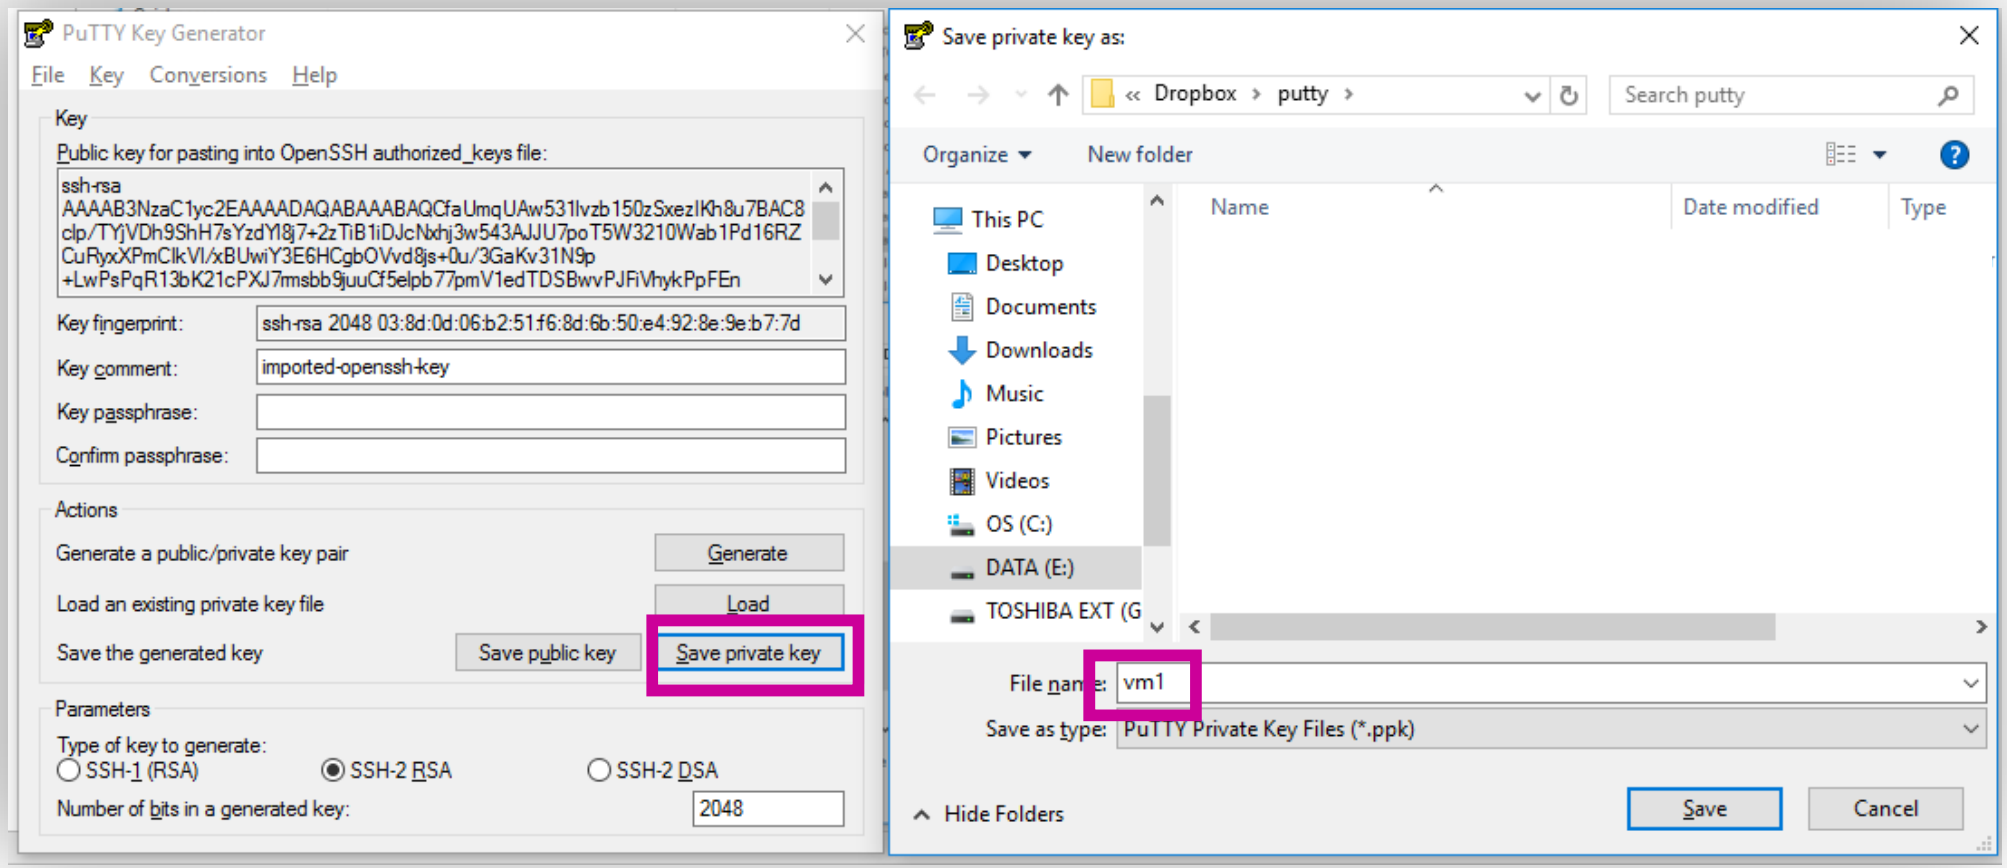

# Allow Network Traffic

Let us try to connect vm1

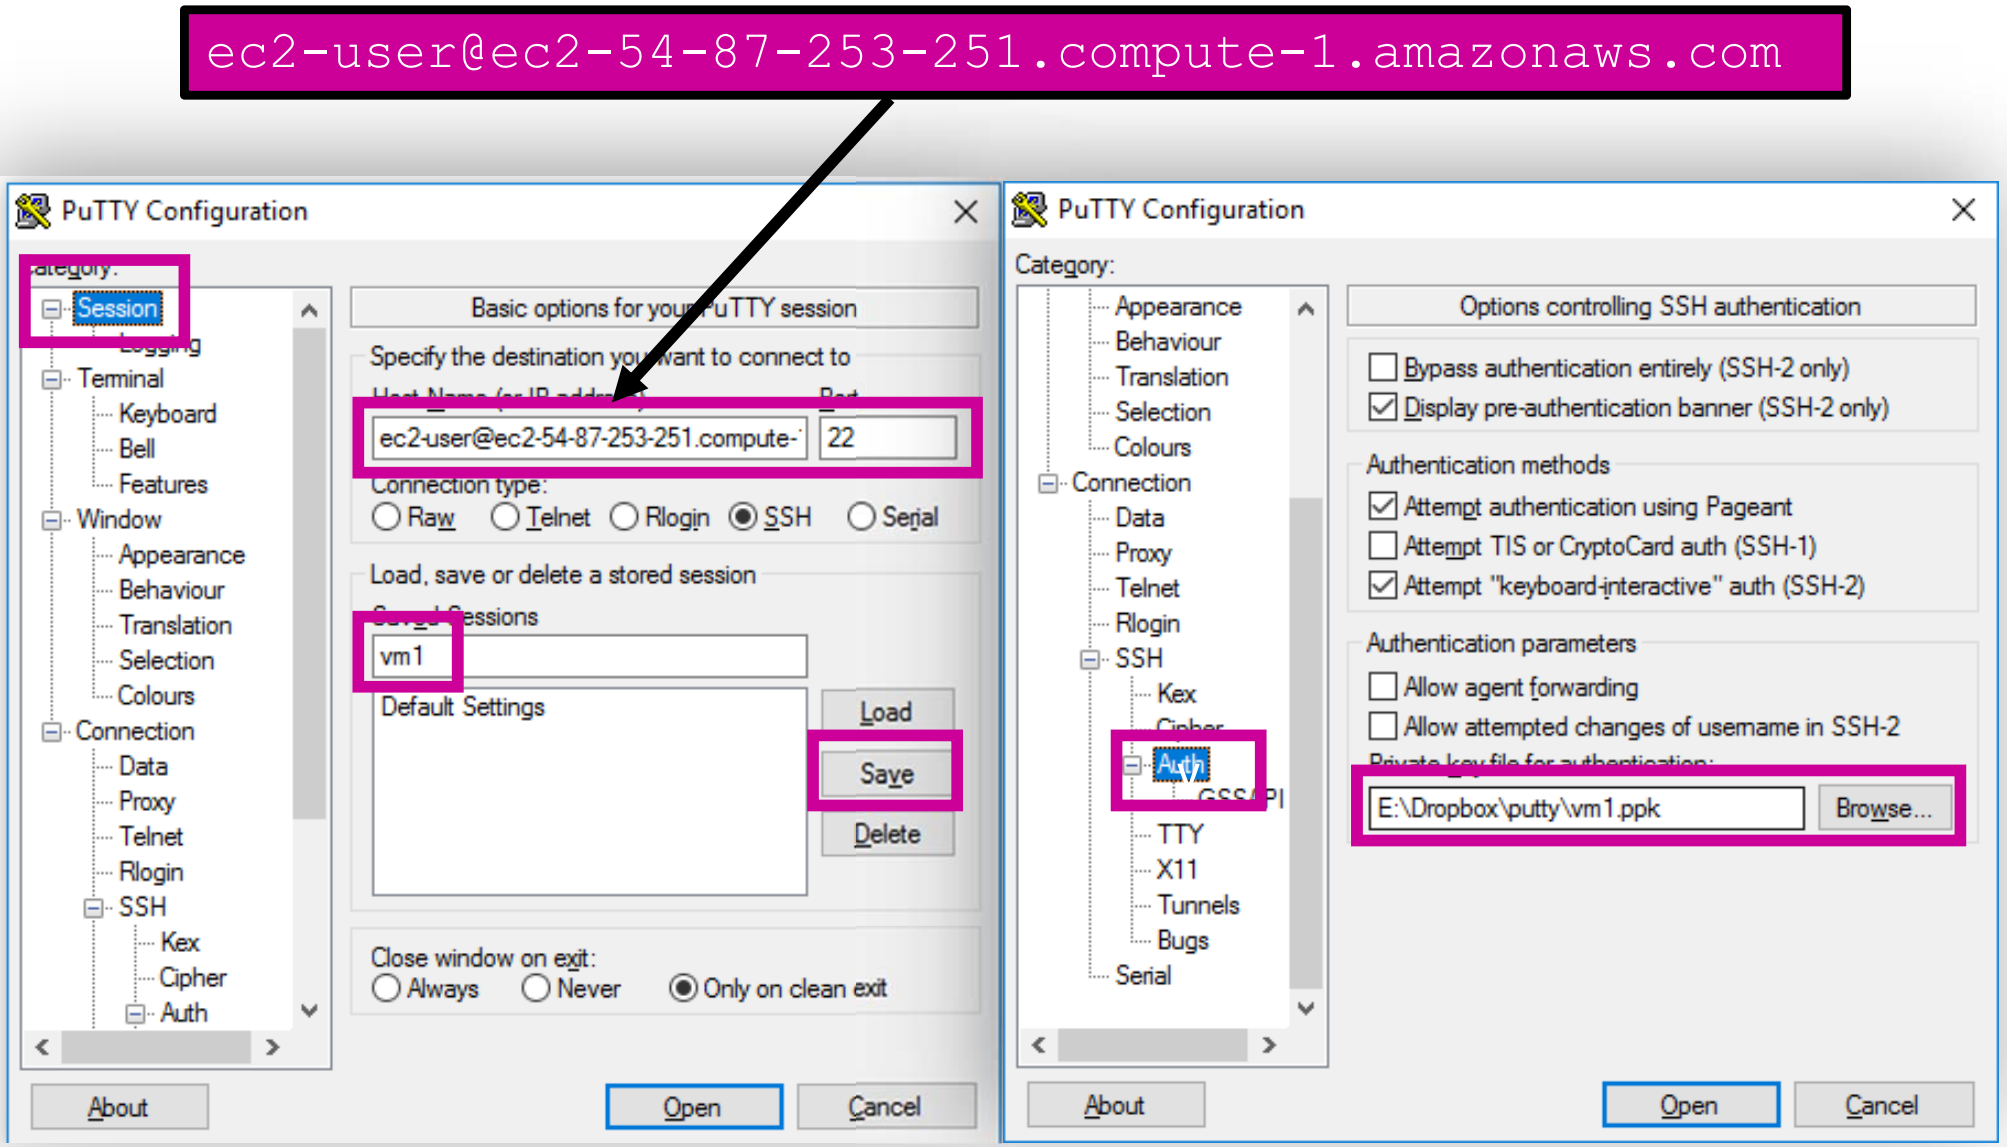

# Allow Network Traffic

Now you have remotely logged into vm1
Let us issue a command **`ifconfig`**

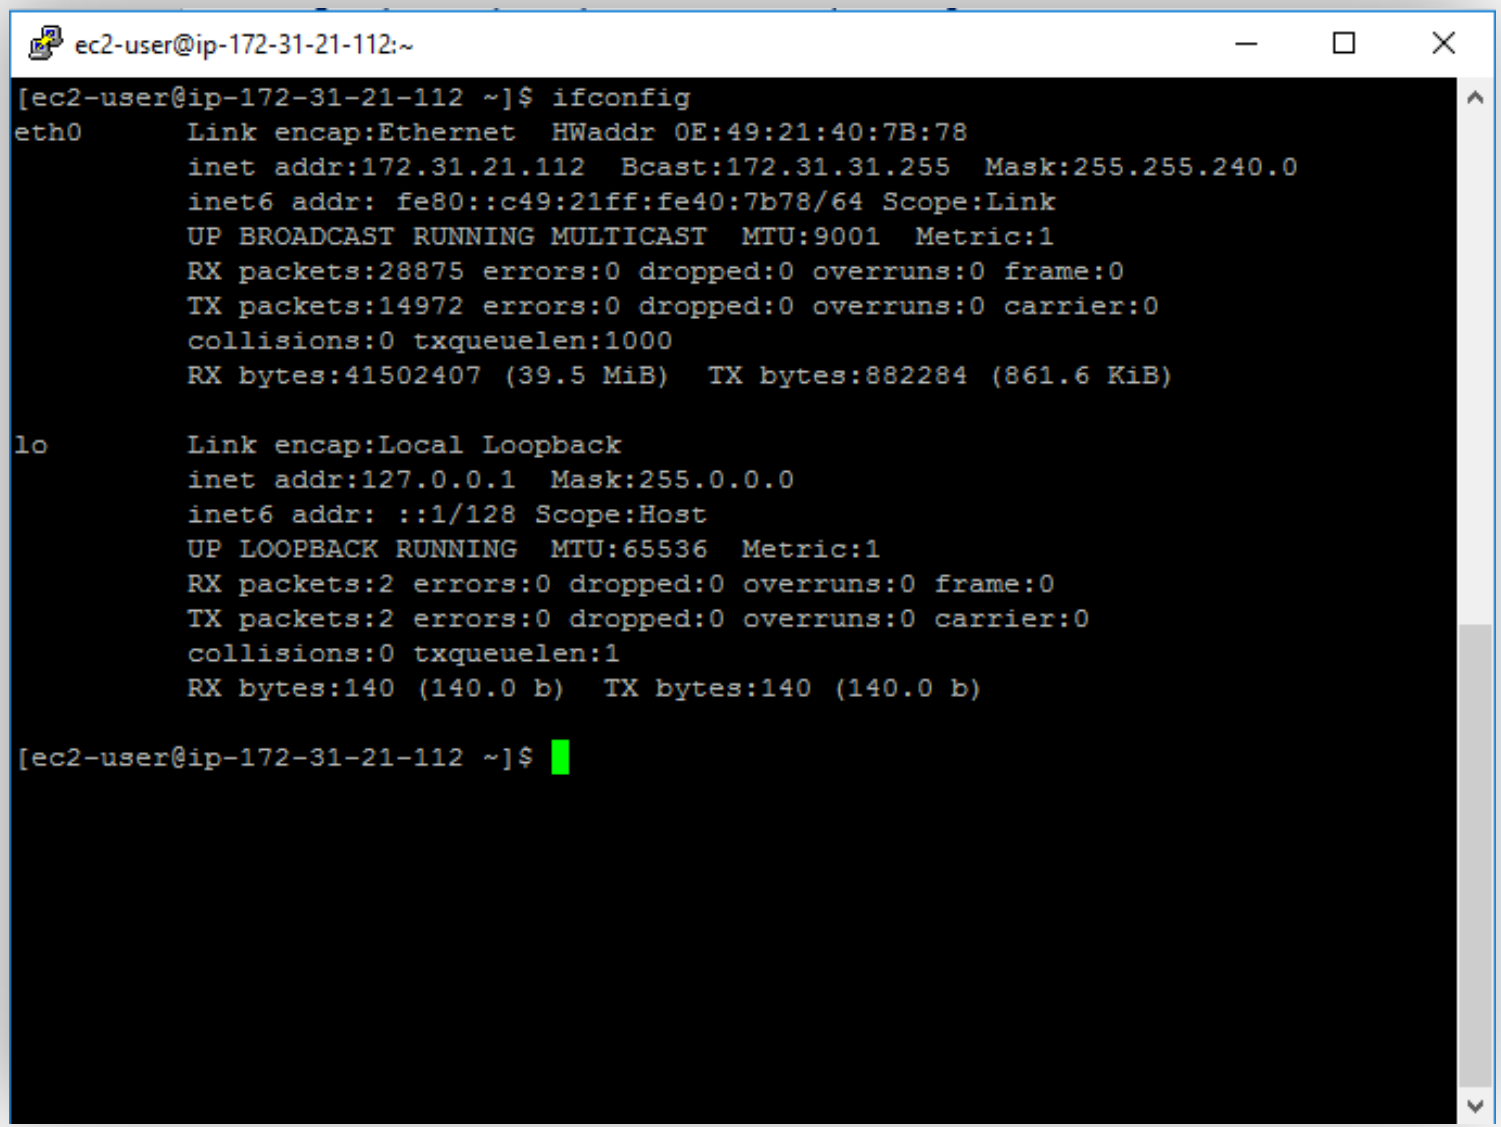

# Allow Network Traffic

Do the same to connect vm2

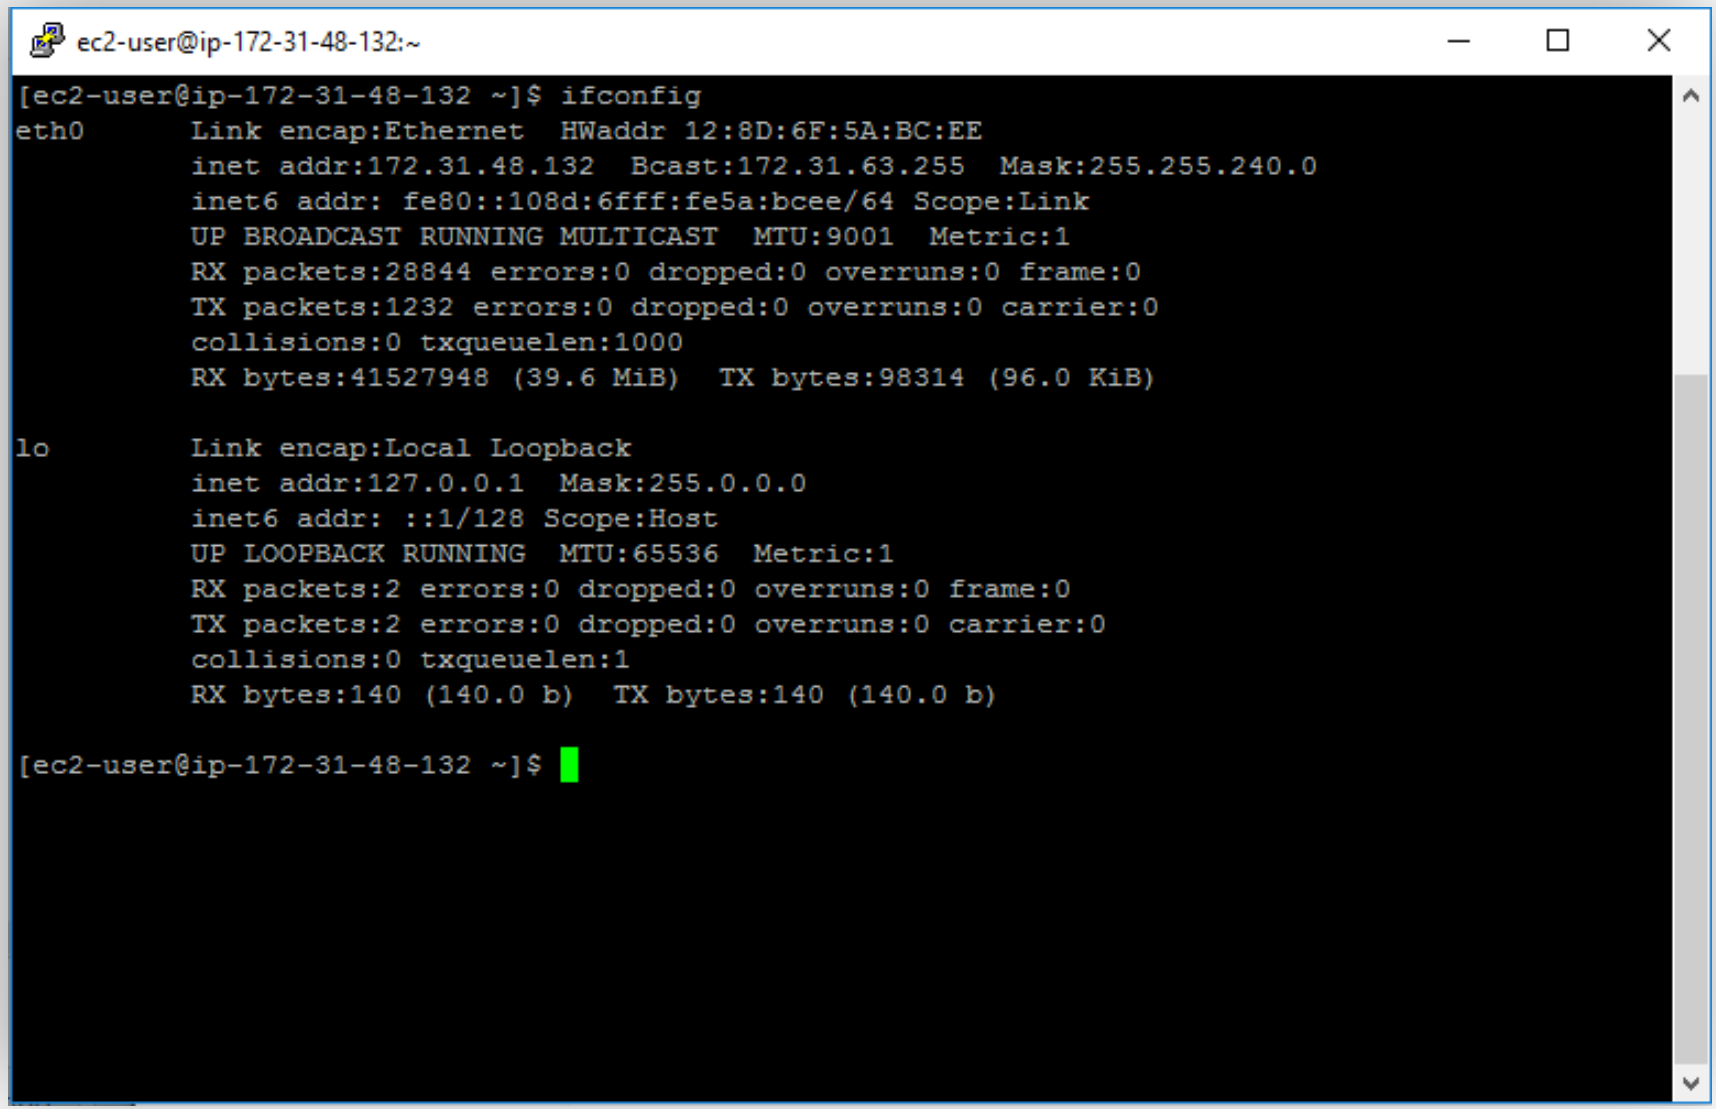

# Allow Network Traffic

Let us see if vm1 and vm2 can talk to each other

```
ping 172.31.48.132
ping 172.31.21.112
```

Cannot ping each other…
Use `ctrl + c` to kill the ping commands

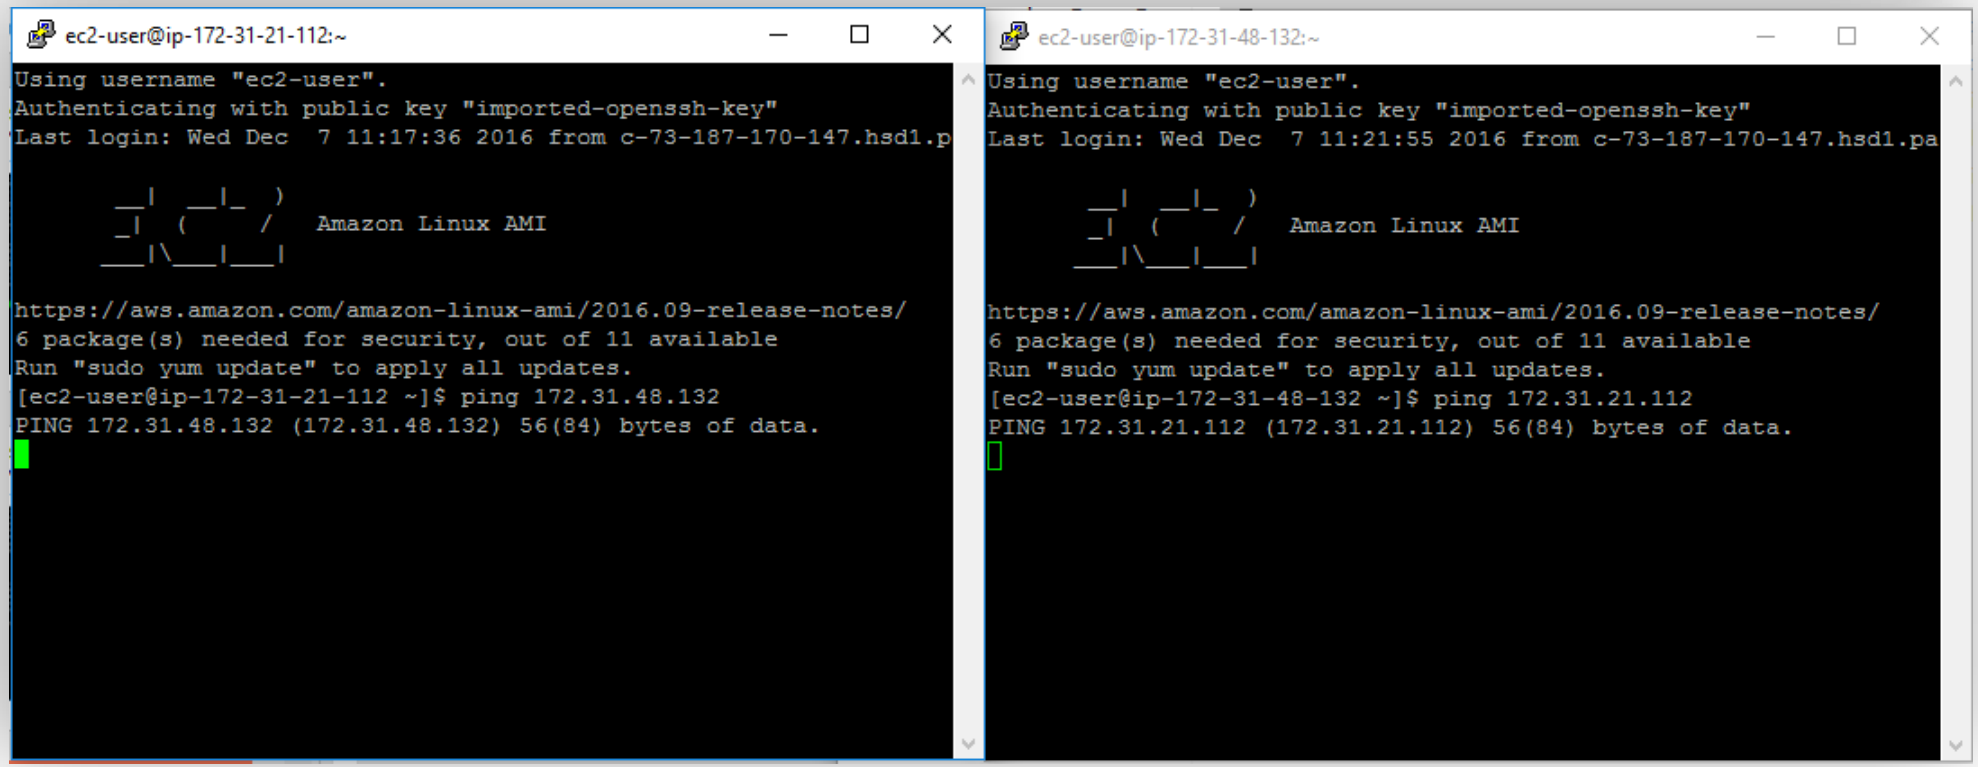

# Allow Network Traffic

This's because you need to add appropriate security group for the VMs to allow network traffic

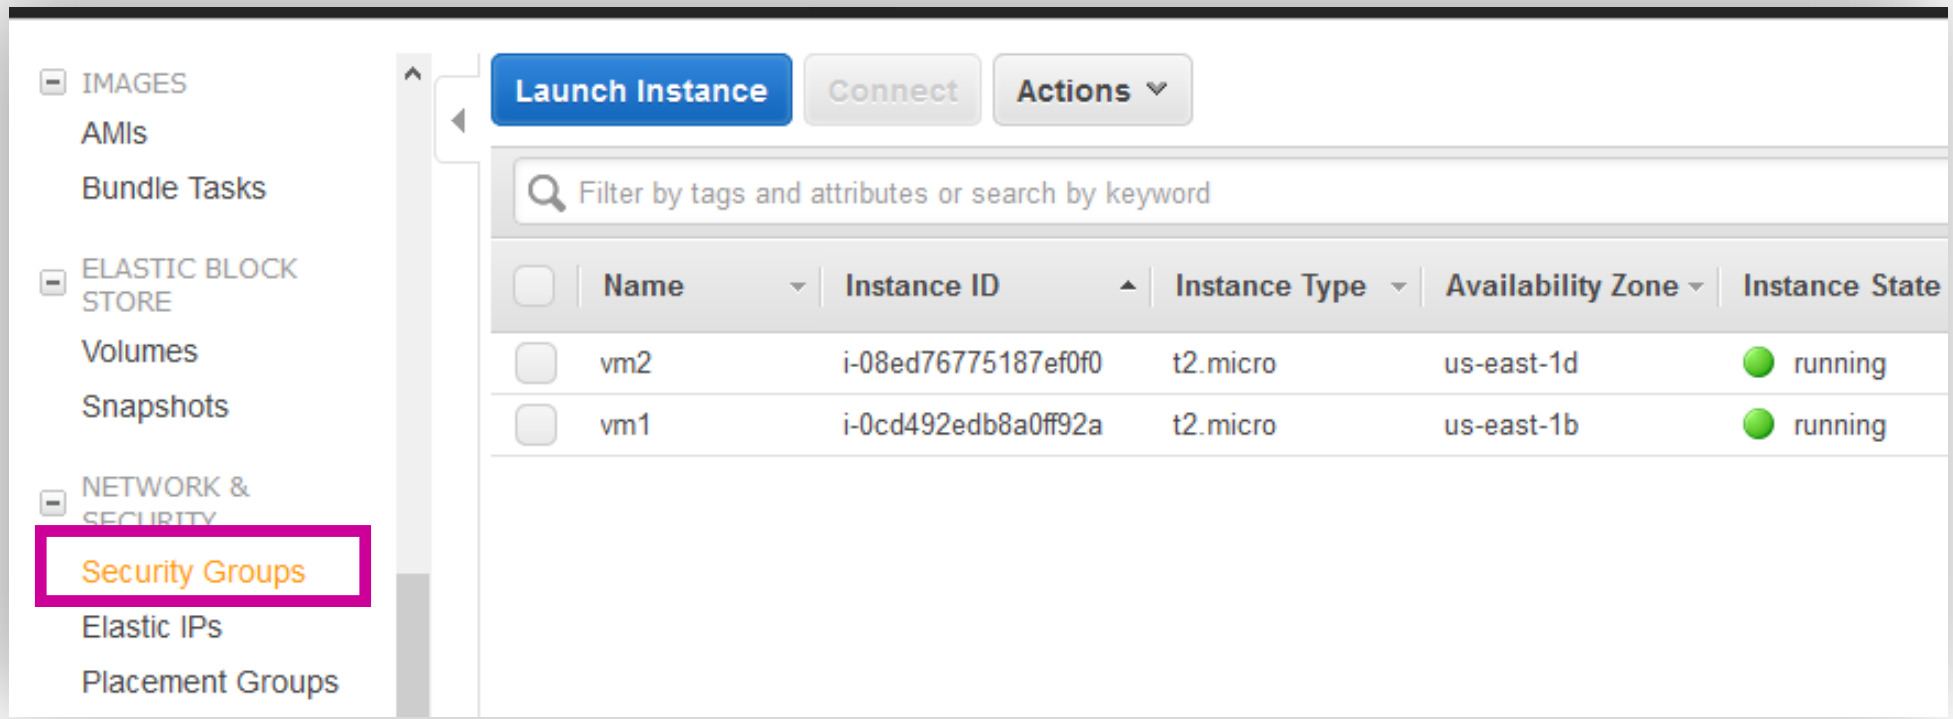

# Allow Network Traffic

Let us allow network traffic for ICMP echo and SSH remote login, and all TCP traffic (not shown)

**Security Group: sg-7e9cd503**

Description | **Inbound** | Outbound | Tags

**Edit**

| Type | Protocol | Port Range | Source |
|---|---|---|---|
| SSH | TCP | 22 | 0.0.0.0/0 |
| All ICMP | All | N/A | 172.31.48.132/32 |

**Security Group: sg-179ed76a**

Description | **Inbound** | Outbound | Tags

**Edit**

| Type | Protocol | Port Range | Source |
|---|---|---|---|
| SSH | TCP | 22 | 0.0.0.0/0 |
| All ICMP | All | N/A | 172.31.21.112/32 |

# Allow Network Traffic

Now you can ping each other

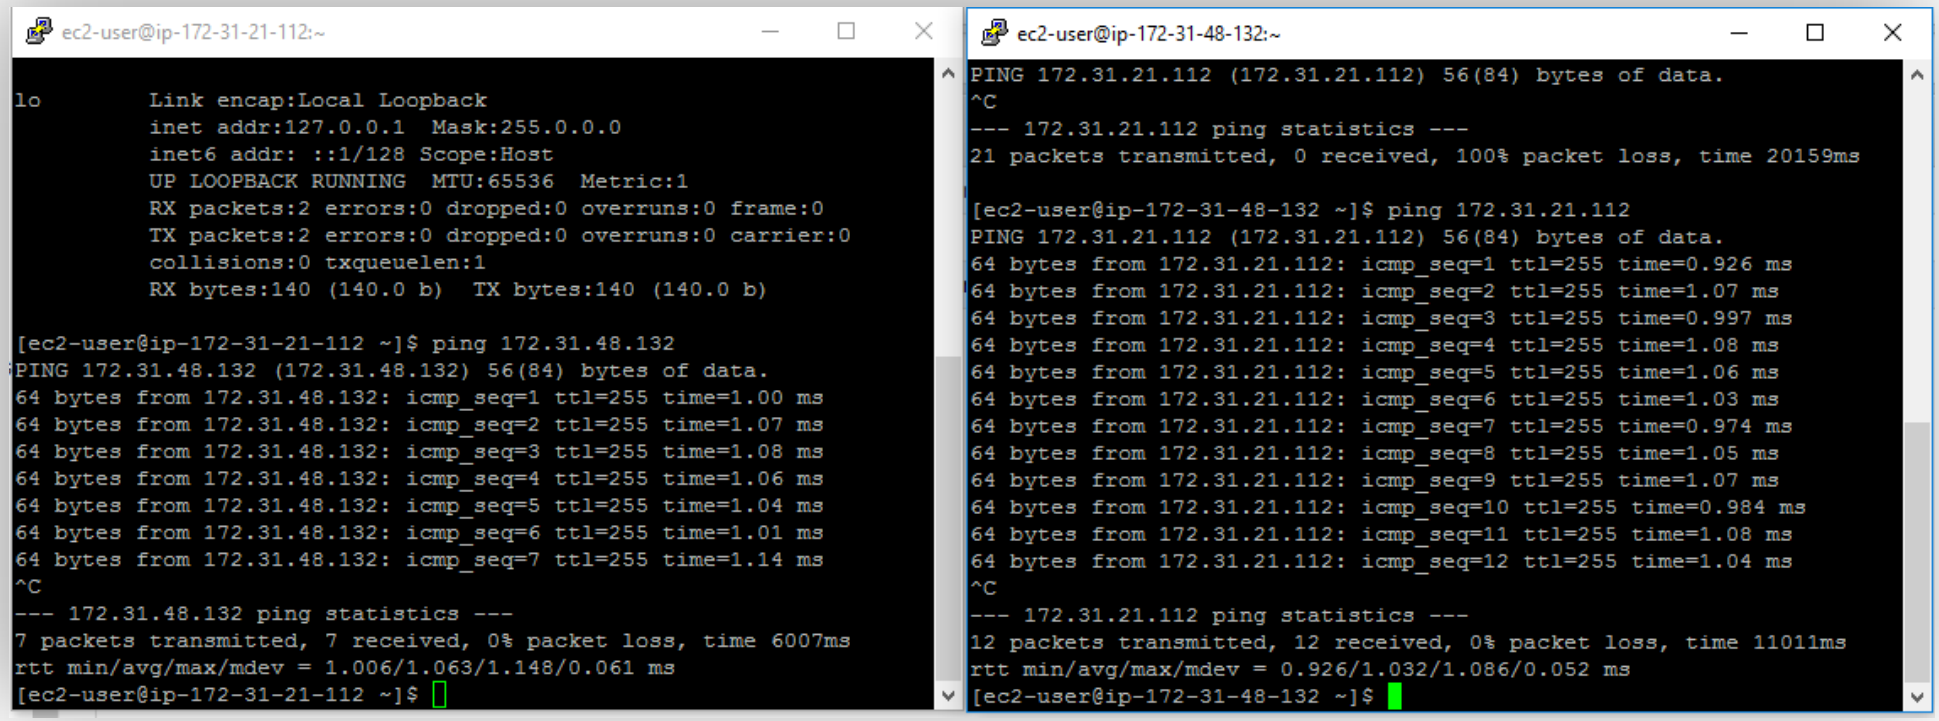

# Transfer Files

Let us using **WinSCP**
https://winscp.net/eng/download.php
https://winscp.net/download/WinSCP-5.9.3-Setup.exe

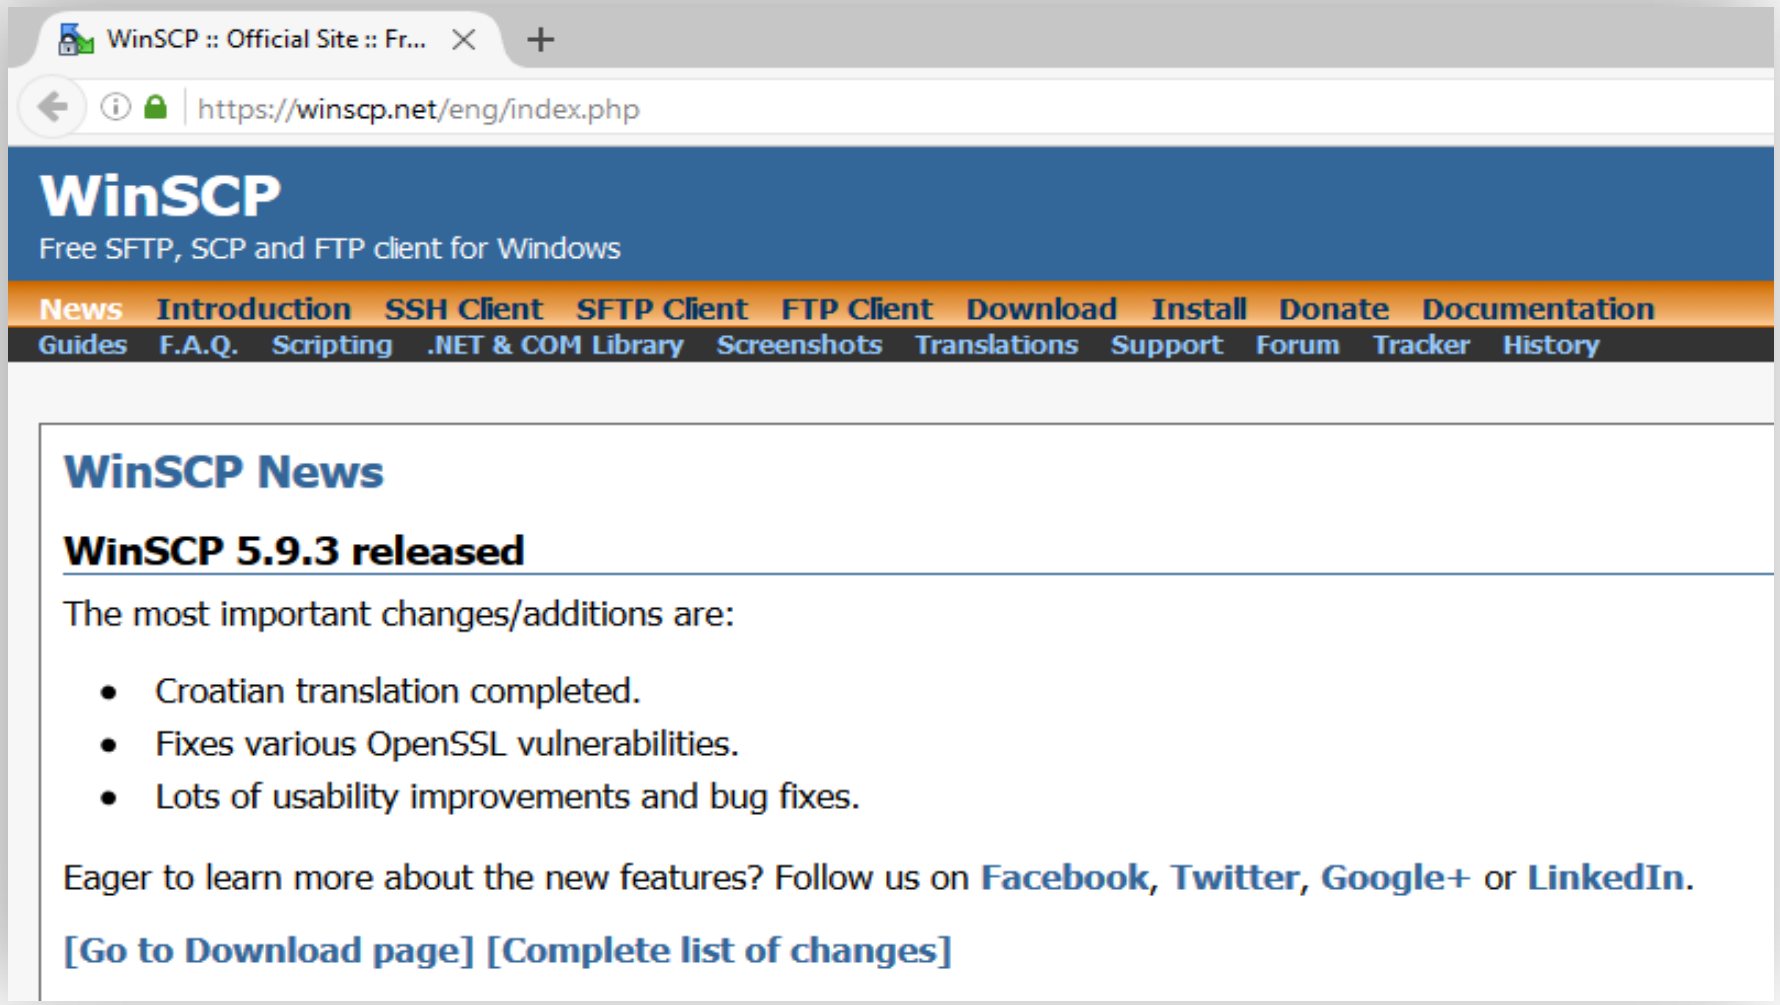

# Transfer Files

## Tools → Import Sites → OK

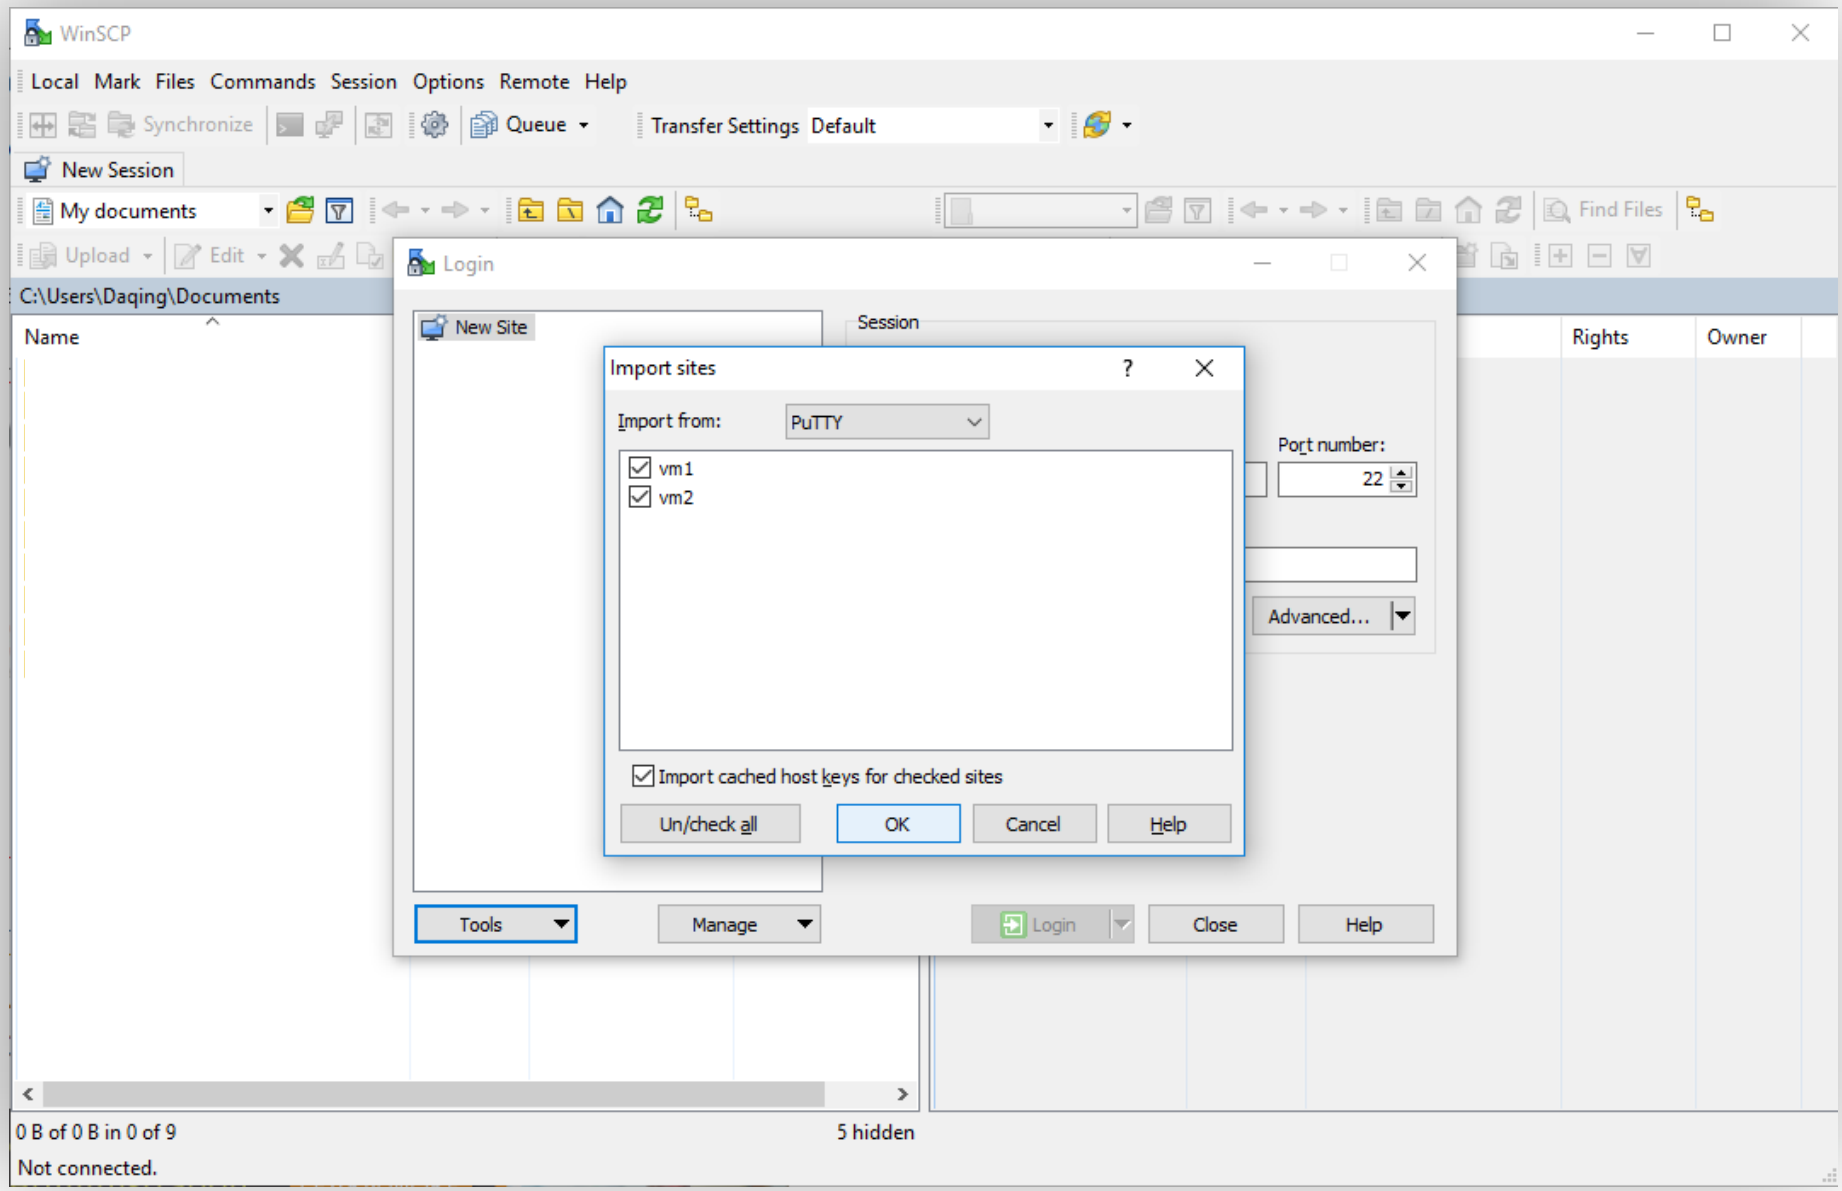

# Transfer Files

Upload the key vm2.pem to VM vm1

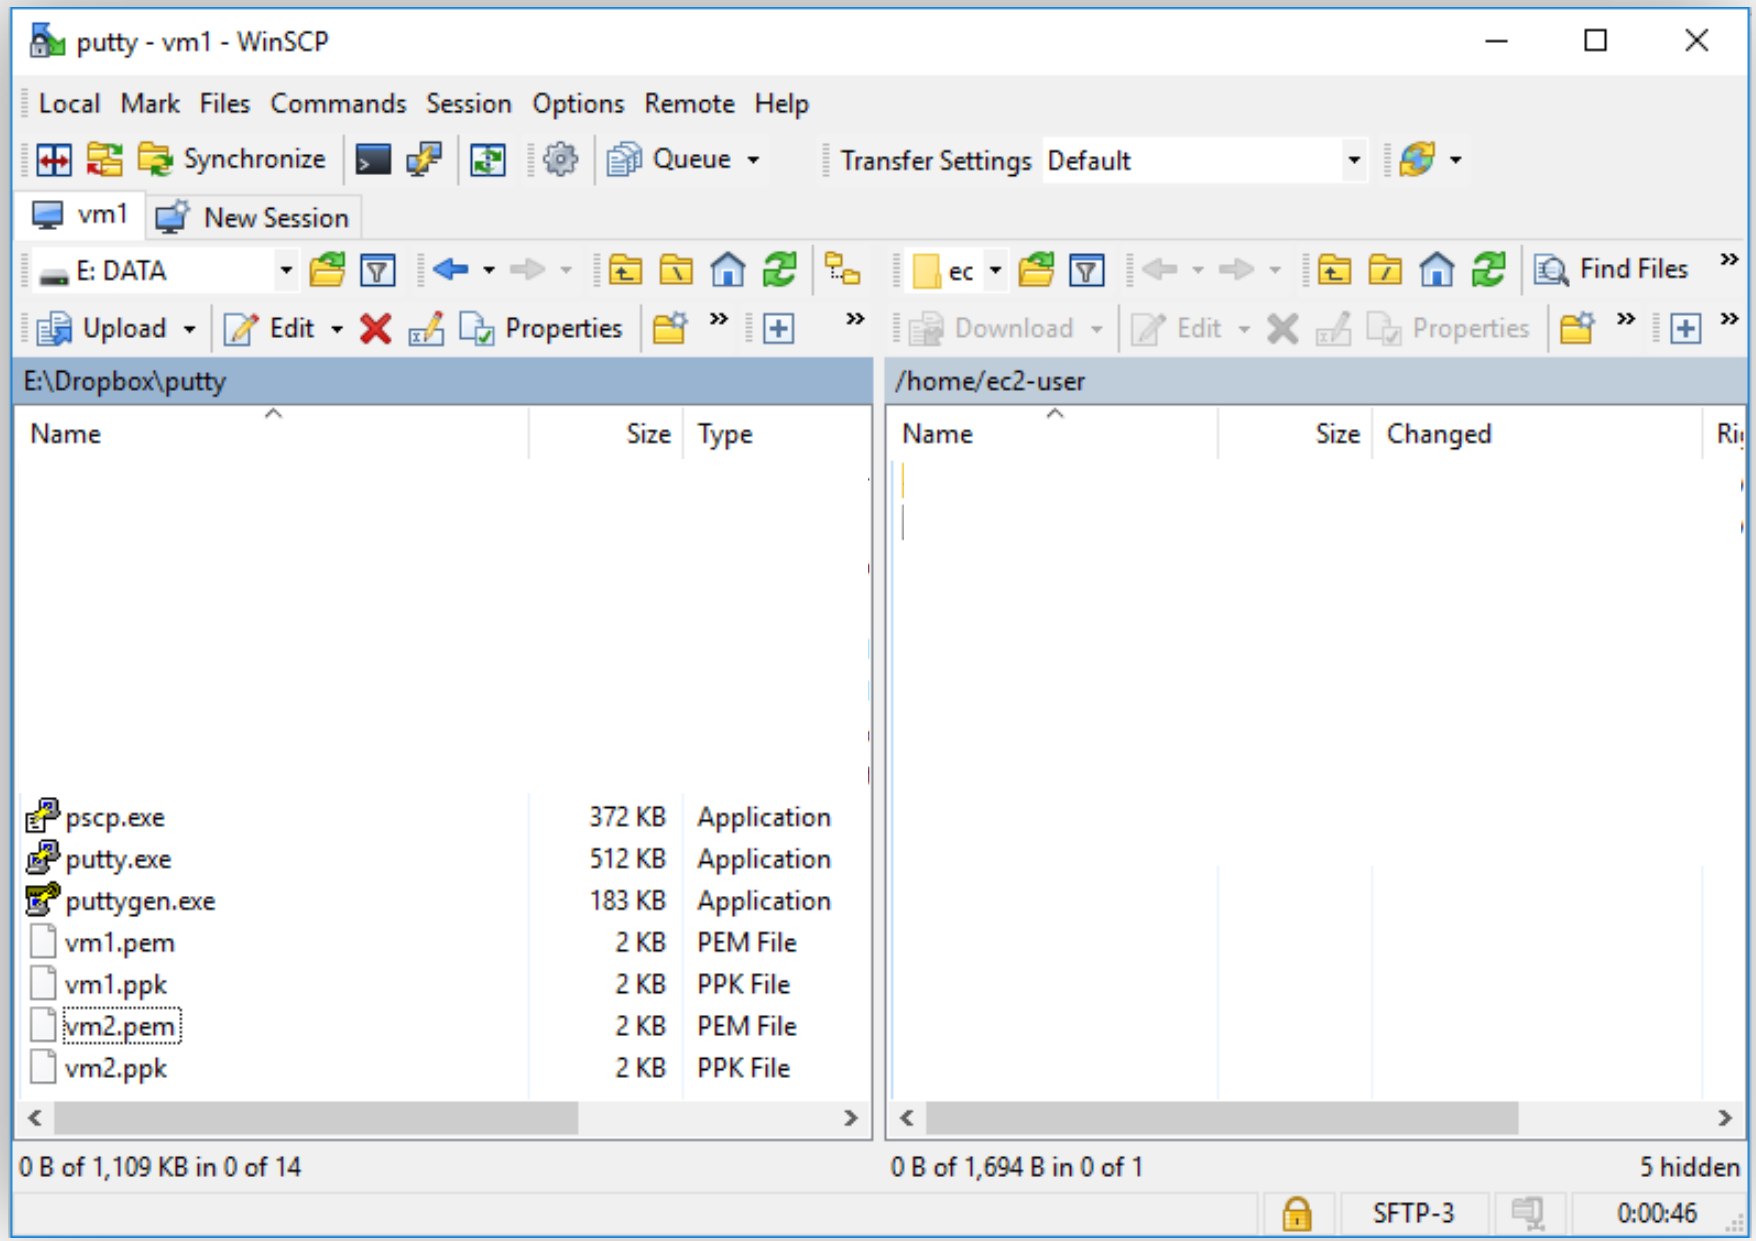

# Remote Access

Log into vm2 from vm1

```
chmod 400 ./vm2.pem
ssh -i ./vm2.pem ec2-user@172.31.48.132
```

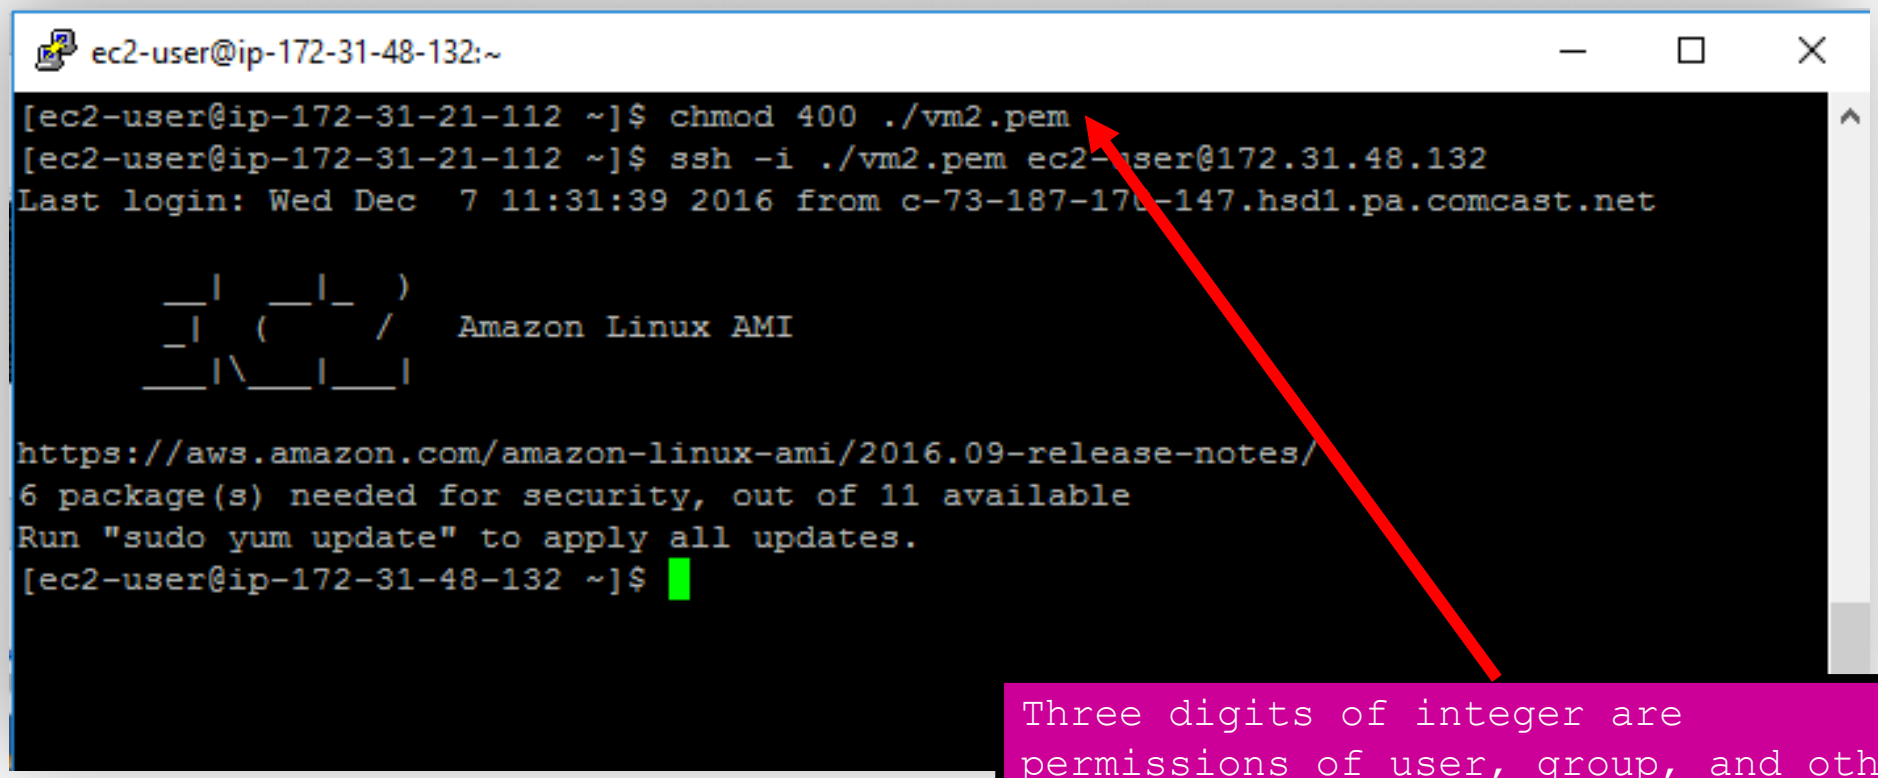


Three digits of integer are permissions of user, group, and others
Each digit can be: 4, 2, 1, 0
4: read; 2: write; 1: execute;
0: no permission
Each digit is a combination of the numbers 4, 2, 1, and 0

# Remote Access

**Log in “without” specifying a key/password**

Remember we’ve uploaded vm2.pem to vm1, and we are able to log into vm2 from vm1

Now at vm1, generate public/private key pairs using **ssh-keygen**

```
[ec2-user@ip-172-31-21-112 ~]$ ssh-keygen
Generating public/private rsa key pair.
Enter file in which to save the key (/home/ec2-user/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/ec2-user/.ssh/id_rsa.
Your public key has been saved in /home/ec2-user/.ssh/id_rsa.pub.
The key fingerprint is:
56:3a:9b:d8:f8:ab:42:01:48:26:2f:62:56:d3:f3:db ec2-user@ip-172-31-21-112
The key's randomart image is:
+--[ RSA 2048]----+
|o+ o.            |
|+.o .o           |
|oo..  o   .      |
|+.  .  . o       |
|     .  S        |
|    .  * E       |
|   .  o +        |
|    .  .         |
|     ...o.       |
+-----------------+
```

Copy vm1’s public key to vm2 using **scp** command

```
[ec2-user@ip-172-31-21-112 ~]$ scp -i ./vm2.pem .ssh/id_rsa.pub ec2-user@172.31.48.132:~/.ssh/
id_rsa.pub                                    100%  407     0.4KB/s   00:00
```

# Remote Access

Now at vm2, put vm1's public key into **authorized_keys**

```
[ec2-user@ip-172-31-48-132 .ssh]$ cat id_rsa.pub >> authorized_keys
[ec2-user@ip-172-31-48-132 .ssh]$ cat authorized_keys
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAABAQCRKKRBGLpZawKQn2o8oSfe2zZuY/Gy+wxlz5SAjy1VCmb6HUJhAWw265C1iu5zx
LQZ6iNHlU+1LnUhIkUmelBhcgtCCA828BDhL4cigrmeOS/6zOUMuFrG3En+29lof/XJrqVSxyG/uUZbuoOlx2KbCUvMgofsCJADBV
LwaNGEiSVbXVwizVX6F+4A5OzenSv4ws6Al5N/KZUXaZ+mYTQfzkVZYTMSXOP2QGADVMbeoJ6Wk1ihbSHySIzy5bpOAvmHqDgDje+
D+lBmchmdN6JWIAtK7HdLFurAqCqVlvG5RjnjjCCrdm3skWa+ERUat3v7ZAvfnW8xAEznyN2Wfi4F vm2
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAABAQDNSVfNfJDmowMKiSaQrJ47R+ZziTxUnC+efBUuZkqLfPtBH/NYTpisRAkonanf3
CYc/R8MIwsmcvQn3TEAdporZPnqFSw940N+o2d8bMi4OZ+ybKsW2d5DIrGQGqJ/3FRq8VSTb47Czzqc74T2soovnXmrvC1icaIn3g
GvYB7EJC7nZ21PxrzcjmXBqHfSCPOPtWk02DejaBip4O+i8hT9BA13dClb88+AUiRI8UUjCiBhz9MuD9oMlilpXuitK/HDUadoMTW
yPXBWMnQ8jPDuUhbeP0cv7TZ2q9nYDUnLcjJOVVWmMVj0Kxz6aYHjHY0A6nW1dXLZ+mXx0mHro+Lh ec2-user@ip-172-31-21-1
12
```

Now at vm1, log into vm2 "without" specifying a key

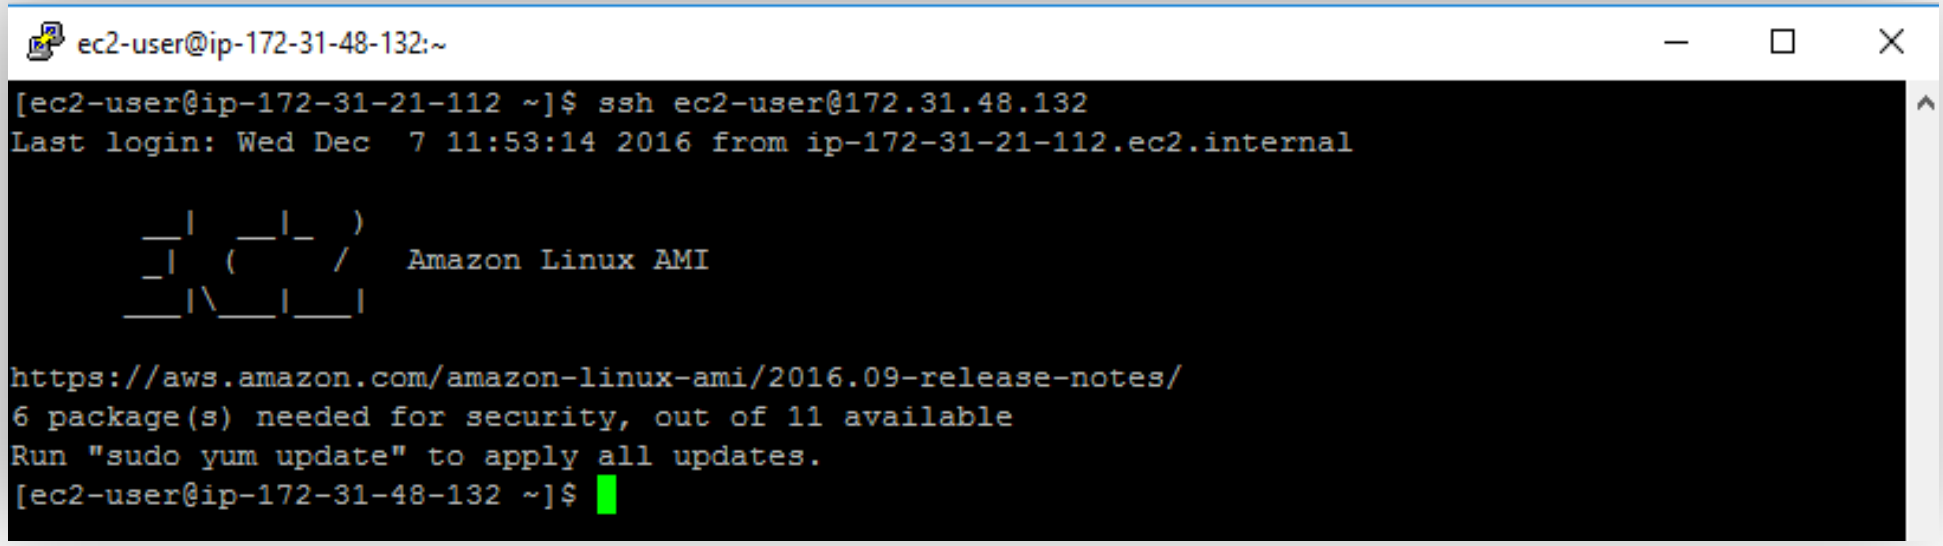

# Remote Access

Do similar things to enable remote log into vm1 from vm2 “without” specifying a key

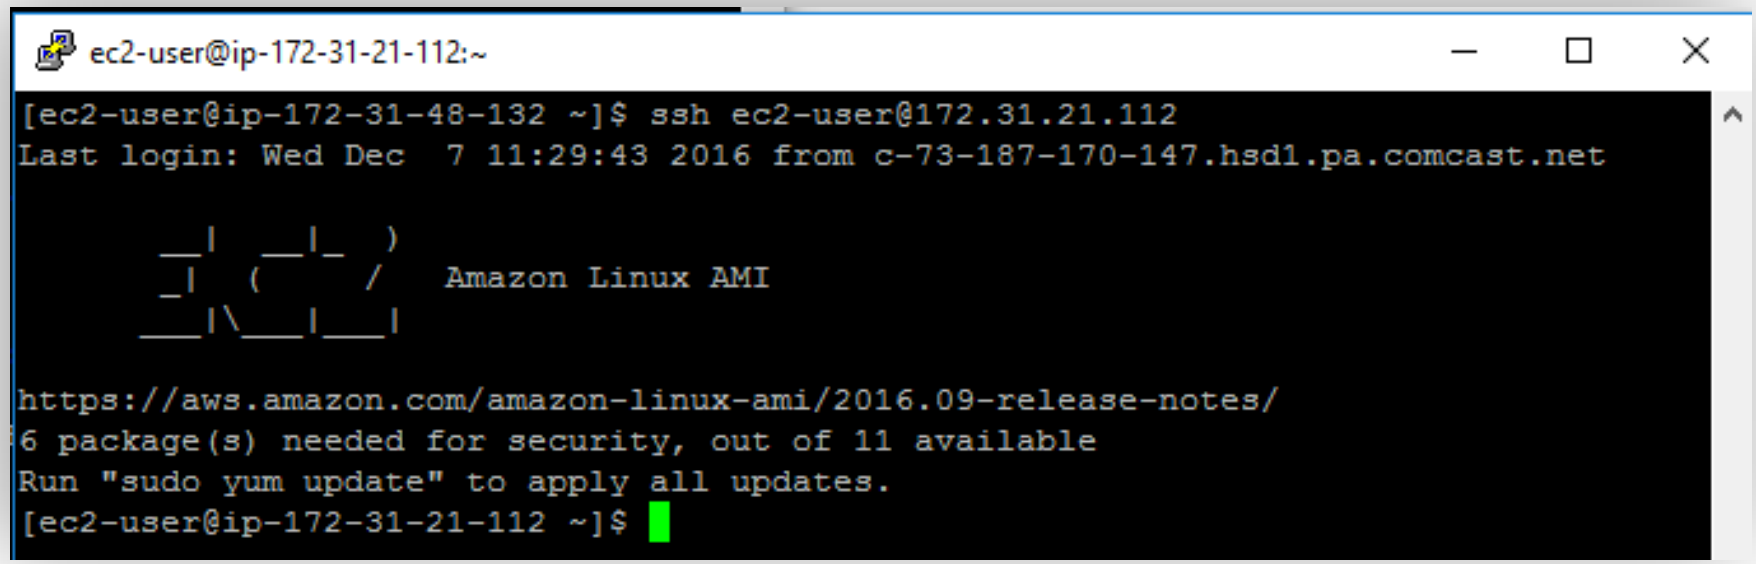

# Hadoop

Download Hadoop, let us download the binary package
http://www-eu.apache.org/dist/hadoop/common/hadoop-2.7.3/hadoop-2.7.3.tar.gz

You can download it into your Windows first, and then transfer it to the virtual machine

Using **wget** command on the virtual machine:

# Hadoop

## Extract the downloaded package, and run it

```
tar -zxvf hadoop-2.7.3.tar.gz
cd hadoop-2.7.3/bin
./hadoop
```

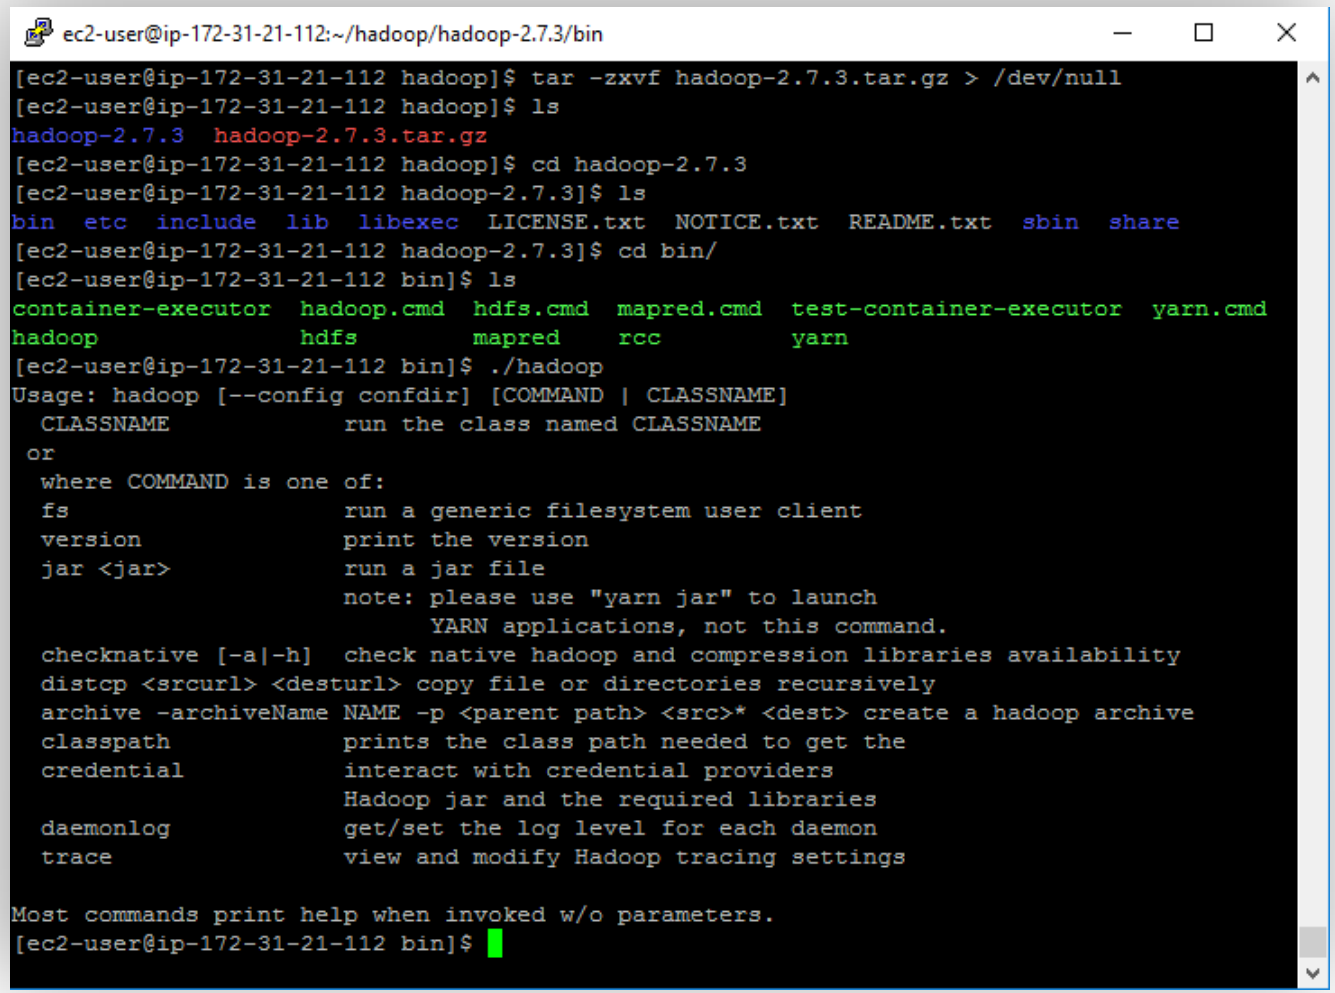

# Hadoop

## Setting up a Single Node Cluster

http://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-common/SingleCluster.html

## Make sure you have JAVA installed

```
env | grep JAVA
```

```
ec2-user@ip-172-31-21-112:~/hadoop/hadoop-2.7.3
[ec2-user@ip-172-31-21-112 hadoop-2.7.3]$ env | grep JAVA
JAVA_HOME=/usr/lib/jvm/jre
[ec2-user@ip-172-31-21-112 hadoop-2.7.3]$
```

## Standalone operation

http://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-common/SingleCluster.html#Standalone_Operation

```
ec2-user@ip-172-31-21-112:~/hadoop/hadoop-2.7.3
[ec2-user@ip-172-31-21-112 hadoop-2.7.3]$ mkdir input
[ec2-user@ip-172-31-21-112 hadoop-2.7.3]$ cp etc/hadoop/*.xml input
[ec2-user@ip-172-31-21-112 hadoop-2.7.3]$ bin/hadoop \
> jar \
> share/hadoop/mapreduce/hadoop-mapreduce-examples-2.7.3.jar \
> grep input output \
> 'dfs[a-z.]+'  /dev/null 2>&1
[ec2-user@ip-172-31-21-112 hadoop-2.7.3]$ cat output/*
1       dfsadmin
[ec2-user@ip-172-31-21-112 hadoop-2.7.3]$
```

**regular expression**